

Core Concepts Section Introduction

Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

Cluster Architecture

1. Kubernetes Architecture
2. ETCD For Beginners
3. ETCD in Kubernetes
4. Kube-api server
5. Controller Manager
6. Kube Scheduler
7. Kubelet
8. Kube-proxy

1. Kubernetes Architecture

For this course we would be using an analogy of Ships to understand the components of a Kubernetes architecture.

Thus Lets consider an example of two ships at sea, one being the cargo ship and the other being the control ship, now the control ship manages the cargo ships, its stipulates which ships goes where and the cargo it should carry, it generally monitors, all the cargo ships running.

In a similar vein relating this to our Kubernetes architecture, we have several nodes (computers) that carry and run our container applications (**worker nodes**), this nodes are controlled by a component of kubernetes called the control plane (**master node**).

Generally the master node, manages, plan, schedule and monitor worker nodes.

Back to our ship analogy, we understand that the control ship is in charge of assigning the loading and unloading of cargo, on and off the cargo ships, specifically assigning which cargo, goes on which ship.

Because of this it is important that we keep track of this information, somewhere.

Relating this back to the kubernetes architecture, **ETCD**, inside the master node is charge of keeping tracking of the information about the containers in key value store. Basically its a database in our kubernetes cluster. To keep track of this information.

Again, Following this pattern, we know that the control ship assigns and loads cargo to a cargo ship based on its Capacity, size e.t.c

For kubernetes, we Load containers onto nodes in similar way, with a component called **kube-scheduler**, kube-scheduler is in charge of loading cargo onto the a particular node, based on the capacity (cpu), memory, tolerance, or any other constrains e.t.c

Now just like a normal ship, we normally have several Departments to deals with several aspect of the ship for example the operations team, is in charge of ship handling, traffic control, e.t.c, while other department such as the cargo ship deals with containers, making sure they are replaced when one is damaged e.t.c

In a similar vein, We have **controllers** in kubernetes that handles different Parts of operations in the kubernetes cluster, the controller manager has components such as node controller, replication controller all handing nodes (if one is down) and getting the desired amount of nodes available (replication) to match the desired state.

One important question to ask thou, is how all these components interact with one another, and who manages them in a broad scope.

This is done by the **kube-api-server**, which manges all these components on a broad level

Now lets Turn our attention back to the cargo ships

Every Cargo ship, has a captain, that monitors activities on each ship, talking to the control ship and getting information e.t.c

In a kubernetes cluster, a **kubelet** is a process that runs on a container and interacts with the kube-api-server, getting information and sending information back

But the kubelet is limited because each one runs on a specific container, what happens when two containers want to communicate with one another. For example in a three tier architecture, with each server either running the client (frontend), backend and the database.

How those the client communicate or connect to the backend and the backend to the database, e.t.c, this communication in a kubernetes cluster is enabled by the **kube-proxy** running on the worker nodes

2. ETC FOR BEGINERS

Objectives

- What is ETCD
- What is key value store
- How to get started quickly
- How to operate ETCD

Later

- What is a distributed System
- How ETCD works
- RAFT Protocol
- Best Practice on number of nodes

What is ETCD



ETCD is a distributed
reliable key-value store
that is Simple, Secure &
Fast

What is key value store

A key value store, is like a relational database where values are stored and retrieved.

Given say, a table with a key called name and a value John Doe, whenever we want to access “John Doe” all we do is get the key name, and the we get “john Doe back”

key-value store

Tabular/Relational Databases

Key	Value
Name	John Doe
Age	45
Location	New York
Salary	5000

Put Name "John Doe"

Get Name

"John Doe"

Name	Age	Location
John Doe	45	New York
Dave Smith	34	New York
Aryan Kumar	10	New York
Lauren Rob	13	Bangalore
Lily Oliver	15	Bangalore



Install ETCD

1. Download Binaries

```
curl -L https://github.com/etcd-io/etcd/releases/download/v3.3.11/etcd-  
v3.3.11-linux-amd64.tar.gz -o etcd-v3.3.11-linux-amd64.tar.gz
```

2. Extract

```
tar xzvf etcd-v3.3.11-linux-amd64.tar.gz
```

3. Run ETCD Service

```
./etcd
```

After installation, And Running the ETCD service

This starts a Process that runs on port **2379**

To save a key in the etcd Database use the command

- `./etcdctl set key1 value1`

To Retrive a value in the etcd Database use the command

- `./etcdctl get key1`

To get all the list of commands use

- `./etcdctl`

Operate ETCD

3. Run ETCD Service

```
./etcd
```

```
▶ ./etcdctl set key1 value1
```

```
▶ ./etcdctl get key1
```

```
value1
```

```
▶ ./etcdctl
```

```
NAME:  
  etcdctl - A simple command line client for etcd.
```

```
COMMANDS:
```

```
  backup      backup an etcd directory  
  cluster-health  check the health of the etcd cluster  
  mk          make a new key with a given value  
  mkdir       make a new directory  
  rm          remove a key or a directory  
  rmdir      removes the key if it is an empty directory or a key-value pair  
  get         retrieve the value of a key
```



3. ETCD in Kubernetes

ETCD In Kubernetes



The ETCD cluster stores several information regarding our kubernetes cluster.

Informations such as the

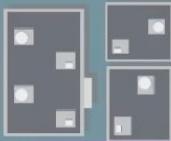
- Pods
- Nodes
- Config
- Secrets
- Accounts
- Roles
- Bindings. e.t.c



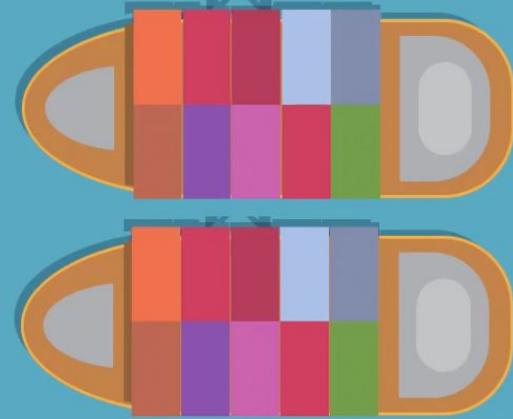
Master

Manage, Plan, Schedule, Monitor
Nodes

- Nodes
- PODs
- Configs
- Secrets
- Accounts
- Roles
- Bindings
- Others



ETCD CLUSTER



There are Two ways to get started with ETCD in kubernetes

1. Is If you are setting up a cluster by you self, in which you would have to set up and install the ECTD binaries your self

Setup - Manual

```
▶ wget -q --https-only \
  "https://github.com/coreos/etcd/releases/download/v3.3.9/etcd-v3.3.9-linux-amd64.tar.gz"
```

etcd.service

```
ExecStart=/usr/local/bin/etcd \
--name ${ETCD_NAME} \
--cert-file=/etc/etcd/kubernetes.pem \
--key-file=/etc/etcd/kubernetes-key.pem \
--peer-cert-file=/etc/etcd/kubernetes.pem \
--peer-key-file=/etc/etcd/kubernetes-key.pem \
--trusted-ca-file=/etc/etcd/ca.pem \
--peer-trusted-ca-file=/etc/etcd/ca.pem \
--peer-client-cert-auth \
--client-cert-auth \
--initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \
--listen-peer-urls https://${INTERNAL_IP}:2380 \
--listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \
--advertise-client-urls https://${INTERNAL_IP}:2379 \
--initial-cluster-token etcd-cluster-0 \
--initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \
--initial-cluster-state new \
--data-dir=/var/lib/etcd
```



2. Is If you are setting up a cluster with Kubeadm, kubeadm install the ECTD binaries in the kube-system namespace

You can access the ectd database with the ectd control utility in this pod, to list all keys in the ectd database, write

```
Kubectl exec etcd-master -n kube-system etcdctl get / --pre
```

I Setup - kubeadm

```
▶ kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcdf6894-prwvl	1/1	Running	0	1h
kube-system	coredns-78fcdf6894-vqd9w	1/1	Running	0	1h
kube-system	etcd-master	1/1	Running	0	1h
kube-system	kube-apiserver-master	1/1	Running	0	1h
kube-system	kube-controller-manager-master	1/1	Running	0	1h
kube-system	kube-proxy-f6k26	1/1	Running	0	1h
kube-system	kube-proxy-hnzsw	1/1	Running	0	1h
kube-system	kube-scheduler-master	1/1	Running	0	1h
kube-system	weave-net-924k8	2/2	Running	1	1h
kube-system	weave-net-hzfcz	2/2	Running	1	1h

Explore ETCD

```
kubectl exec etcd-master -n kube-system etcdctl get / --prefix -keys-only  
  
/registry/apiregistration.k8s.io/apiservices/v1.  
/registry/apiregistration.k8s.io/apiservices/v1.apps  
/registry/apiregistration.k8s.io/apiservices/v1.authentication.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.autoscaling  
/registry/apiregistration.k8s.io/apiservices/v1.batch  
/registry/apiregistration.k8s.io/apiservices/v1.networking.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.rbac.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.storage.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1beta1.admissionregistration.k8s.io
```

Run inside the etcd-master POD

Registry

minions

pods

replicasets

deployments

roles

secrets



In a highly available Environment ECTD, where they are multiple master nodes in the kubernates cluster, then we will have multiple ETCD instances in the Kubernetes cluster.

ETCD in HA Environment



```
ExecStart=/usr/local/bin/etcd \
--name ${ETCD_NAME} \
--cert-file=/etc/etcd/kubernetes.pem \
--key-file=/etc/etcd/kubernetes-key.pem \
--peer-cert-file=/etc/etcd/kubernetes.pem \
--peer-key-file=/etc/etcd/kubernetes-key.pem \
--trusted-ca-file=/etc/etcd/ca.pem \
--peer-trusted-ca-file=/etc/etcd/ca.pem \
--peer-client-cert-auth \
--client-cert-auth \
--initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \
--listen-peer-urls https://${INTERNAL_IP}:2380 \
--listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \
--advertise-client-urls https://${INTERNAL_IP}:2379 \
--initial-cluster-token etcd-cluster-0 \
--initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \
--initial-cluster-state new \
--data-dir=/var/lib/etcd
```



ETCD – Commands (Optional)

(Optional) Additional information about ETCDCTL Utility
ETCDCTL is the CLI tool used to interact with ETCD. ETCDCTL can interact with ETCD Server using 2 API versions – Version 2 and Version 3. By default it's set to use Version 2. Each version has different sets of commands

For example, ETCDCTL version 2 supports the following commands:

```
etcdctl backup
```

```
etcdctl cluster-health
```

```
etcdctl mk
```

```
etcdctl mkdir
```

```
etcdctl set
```

Whereas the commands are different in version 3

```
etcdctl snapshot save
```

```
etcdctl endpoint health
```

```
etcdctl get
```

```
etcdctl put
```

ETCD In Kubernetes – Commands (Optional)

Get etcd installed in kubernetes

- Kubectl get pods -n kube-system
- Or
- Kubectl get pods --namespace=kube-system

Where -n is namespace

To set the right version of API set the environment variable ETCDCTL_API command

```
export ETCDCTL_API=3
```

When the API version is not set, it is assumed to be set to version 2. And version 3 commands listed above don't work. When API version is set to version 3, version 2 commands listed above don't work.

Apart from that, you must also specify the path to certificate files so that ETCDCTL can authenticate to the ETCD API Server. The certificate files are available in the etcd-master at the following path. We discuss more about certificates in the security section of this course. So don't worry if this looks complex:

```
--cacert /etc/kubernetes/pki/etcd/ca.crt  
--cert /etc/kubernetes/pki/etcd/server.crt  
--key /etc/kubernetes/pki/etcd/server.key
```

So for the commands, I showed in the previous video to work you must specify the ETCDCTL API version and path to certificate files. Below is the final form:

```
kubectl exec etcd-controlplane -n kube-system -- sh -c "ETCDCTL_API=3 etcdctl get /  
--prefix --keys-only --limit=10 --cacert /etc/kubernetes/pki/etcd/ca.crt --cert  
/etc/kubernetes/pki/etcd/server.crt --key /etc/kubernetes/pki/etcd/server.key"
```

4. Kube-api server

kube-api
server

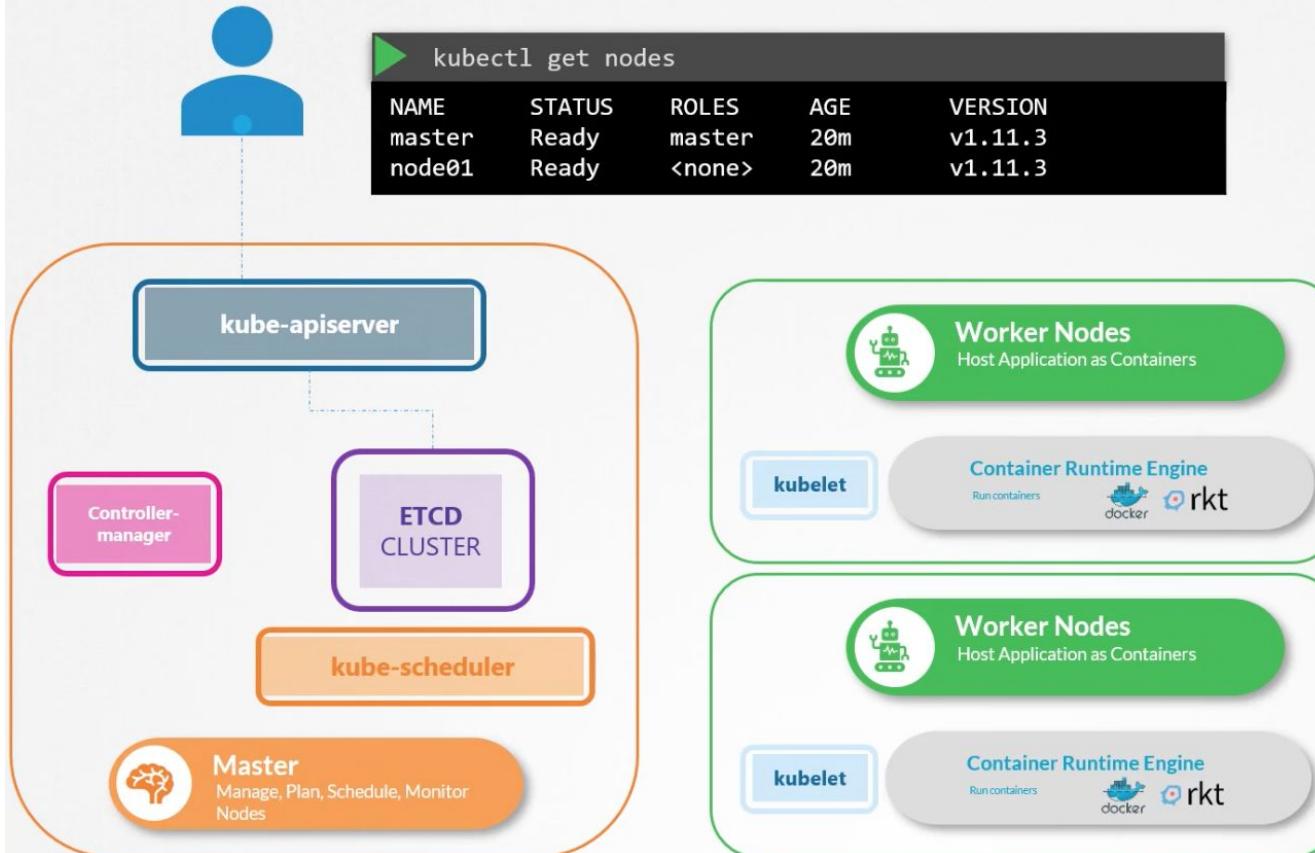


Like we discussed the kube-api-server, is the Primary Management Component in kubernetes, it is the entry point to the kubernetes cluster

For example when we make an kubectl command to get information from the cluster

The Kube-api-server, gets this information from ETCD, and sends a response back to the client

Kubernetes Architecture



The kube-api-server also performs the following functions

I Kube-api Server

1. Authenticate User

2. Validate Request

3. Retrieve data

4. Update ETCD

5. Scheduler

6. Kubelet



Kube Controller Manager

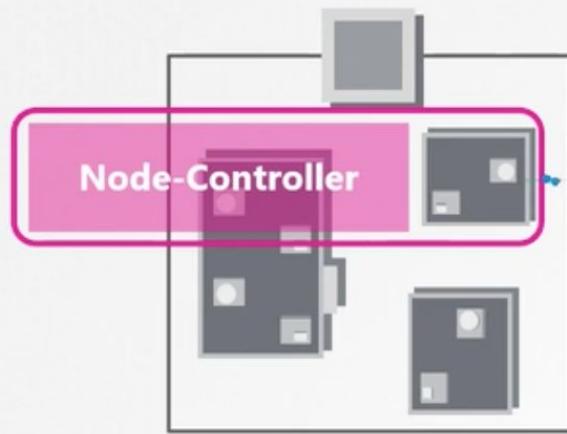


4 - Kube Controller Manager: is a process that manages and monitors the state of various components in the system, and works on bringing the entire system to the desired stated state

For Example the node controller, manges and monitors the nodes in a kubernetes cluster, performing health checks every 5 seconds to make sure, the node is Reachable

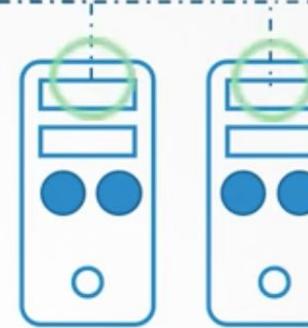
If it finds a Node unreachable, it terms it unhealthy.

Controller



```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0



Watch Status

Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 4us



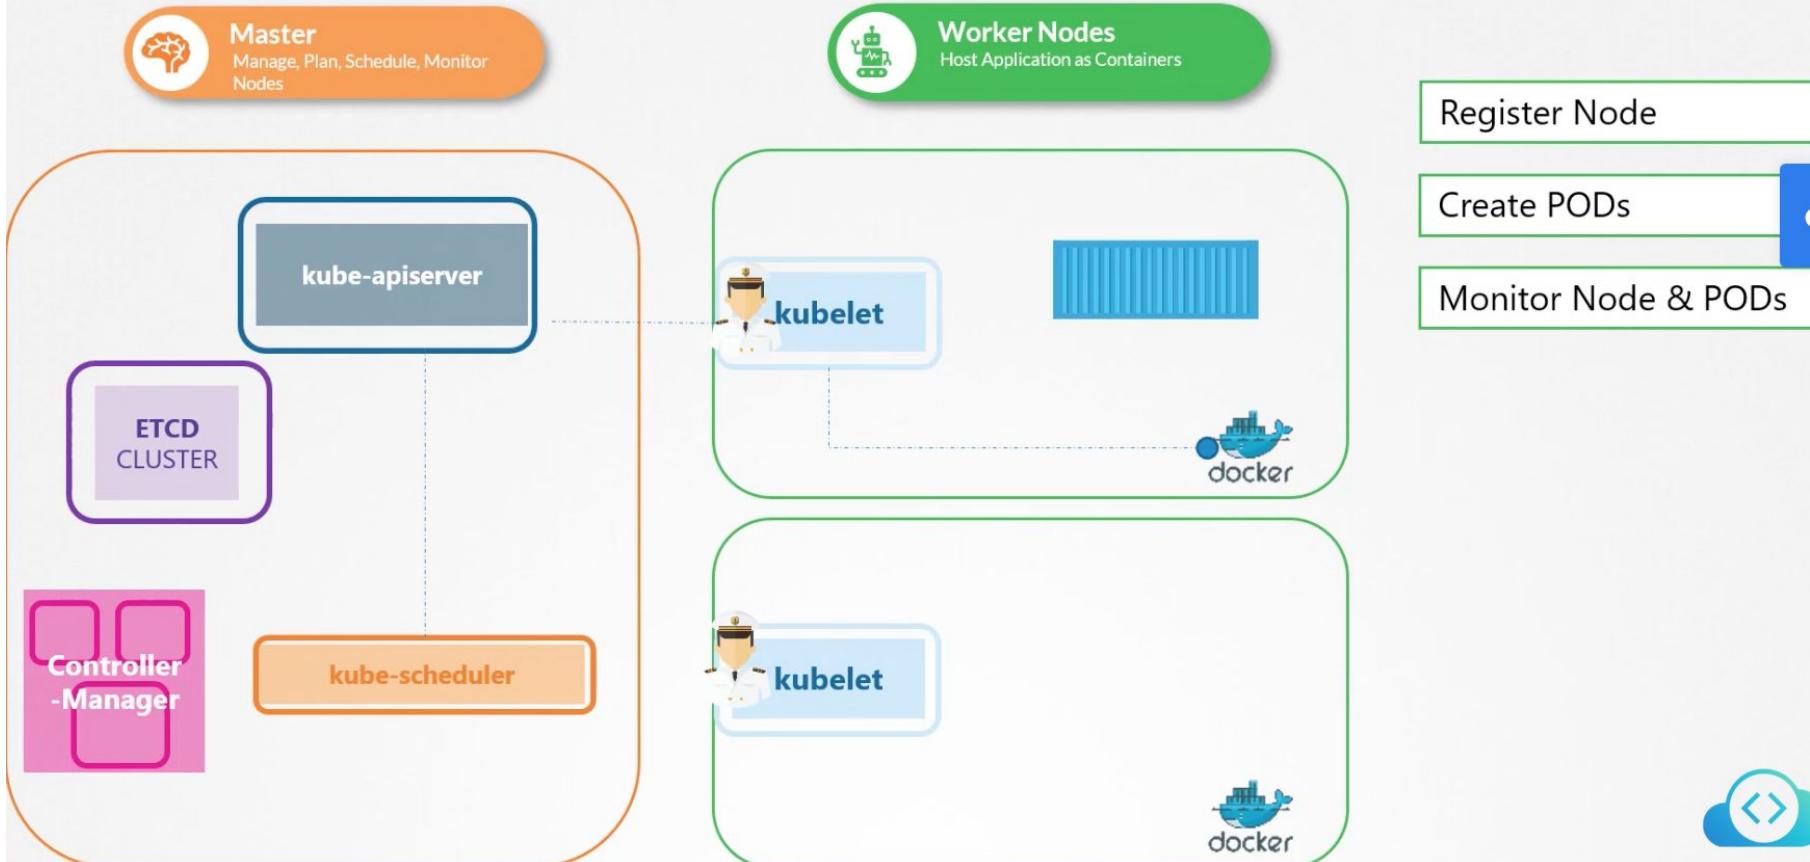
5. Kubelet

Kubelet



As stated Earlier the kubelet, is like the captain of the ship, he in charge of communicating with the API server, sending and receiving Information, from the kube-api-server. E.T.C

Kubernetes Architecture

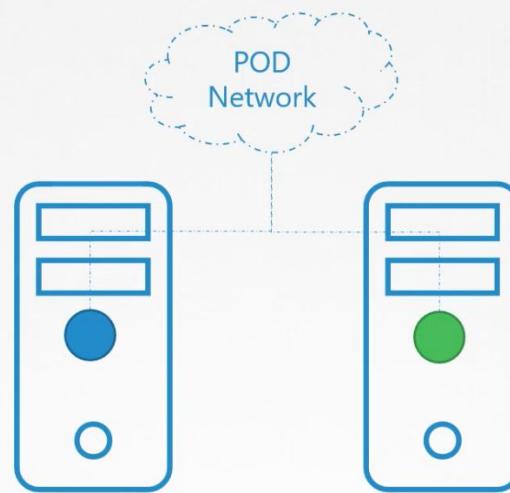


Kube-proxy

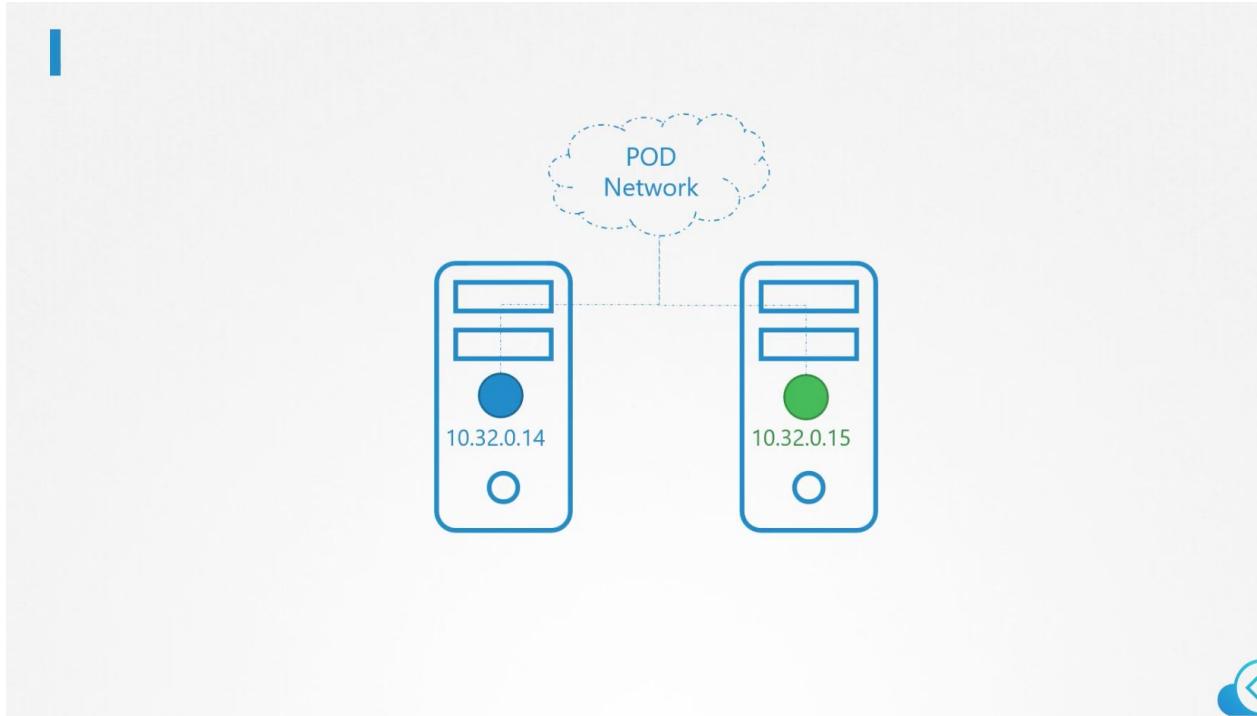


Within a kubernetes cluster, Every Pod can reach another Pod, but how dose it achieve this?

This is done, via a virtual Network existing in the cluster i.e (Pod Network)

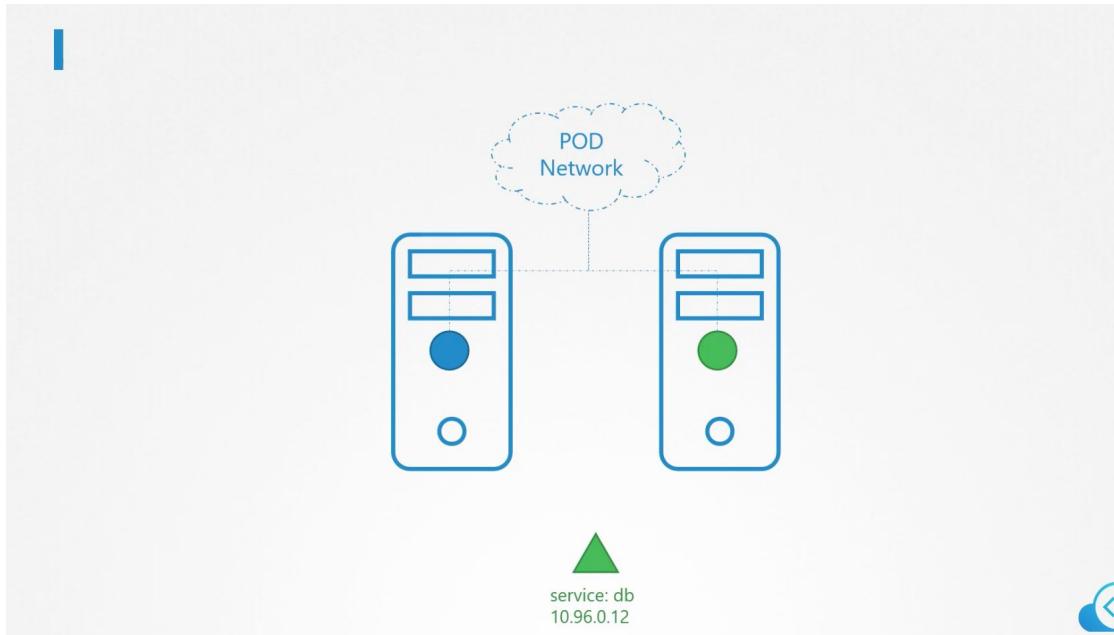


Let say For example in the First Pod, am having a web application running on that server (pod), and on the Second Pod, a database Application, the first Pod(web application) can connect to the second pod (Database application) through its IP address



But because Pods are liable to Fail and be Replaced. This IP Address might not be the best way to connect to an application, cause when it gets replaced, we wont be connecting to the database via the same IP anymore.

A way to solve this problem is to use a service

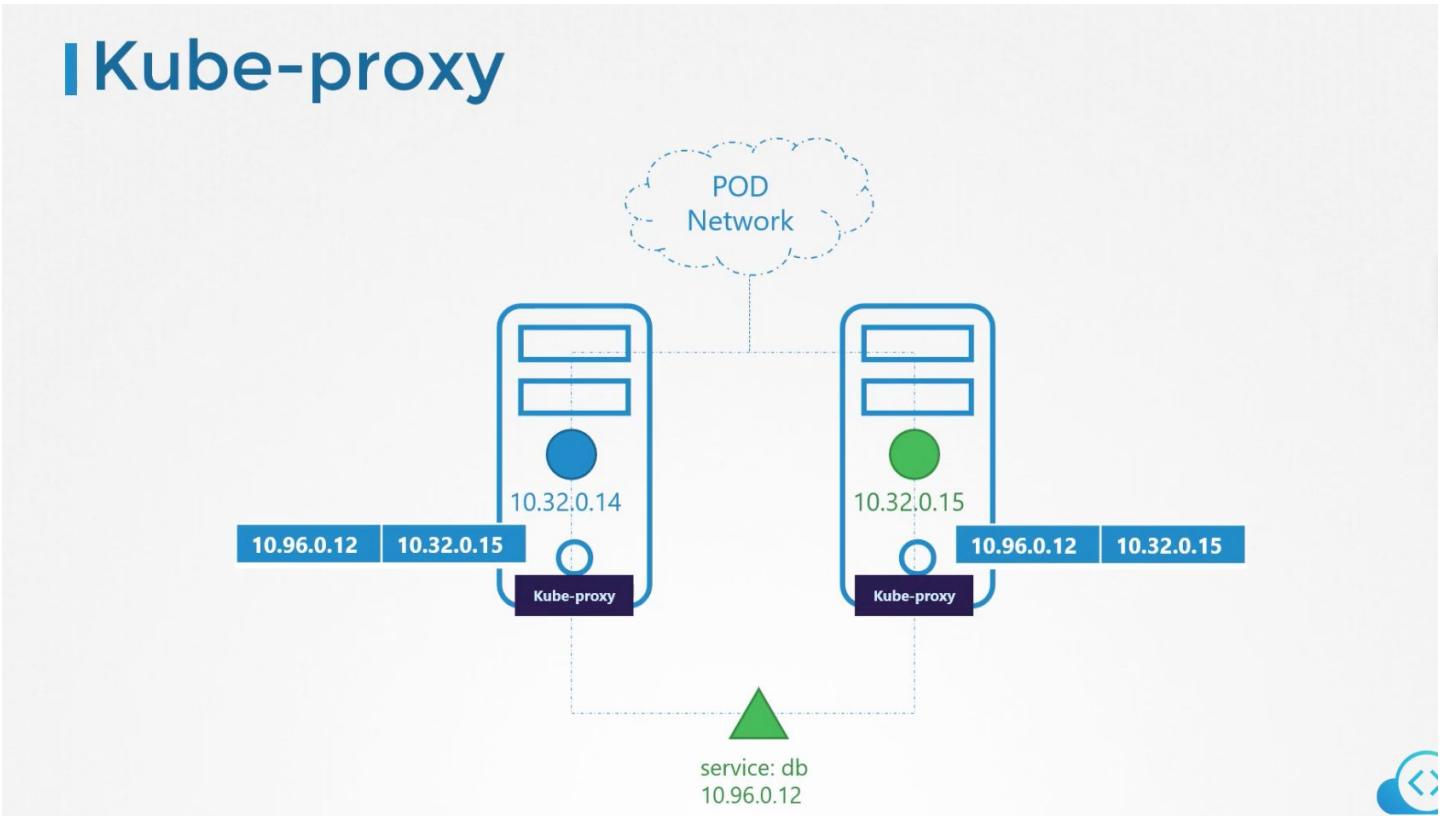


Now the Application can connect to another Pod via the service

The service is simple an abstraction that always has access to the access Pods IP address, even when a Pod is Replaced or changed

Also important to note that the service is an abstract component, that just lives in the clusters Memory, its not like a Pod that we can actually access or share, how do we then make the service available across the cluster

This is achieved with the kube-proxy component, which is a process that runs on a worker node, and always checks and access these services



PODS

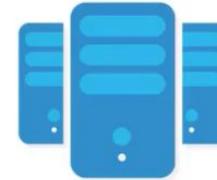
POD



Assumptions



Docker Image



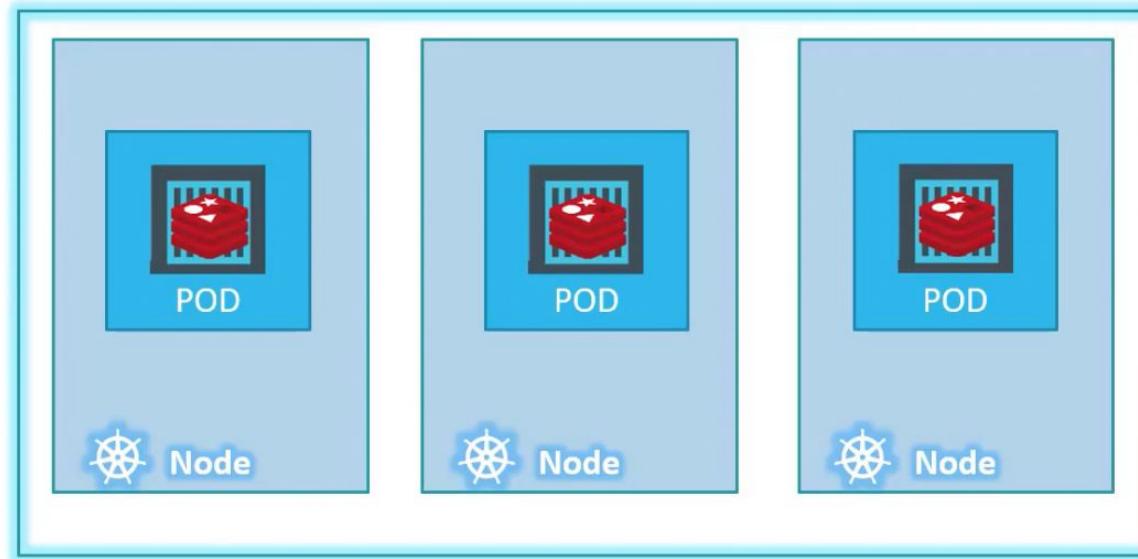
Kubernetes Cluster

A Pod is the smallest unit in a Kubernetes cluster.

Our Goal with kubernetes is to deploy our applications as containers, but in kubernetes we usually interact with pods

And these Pods runs a containers, which is what runs our application

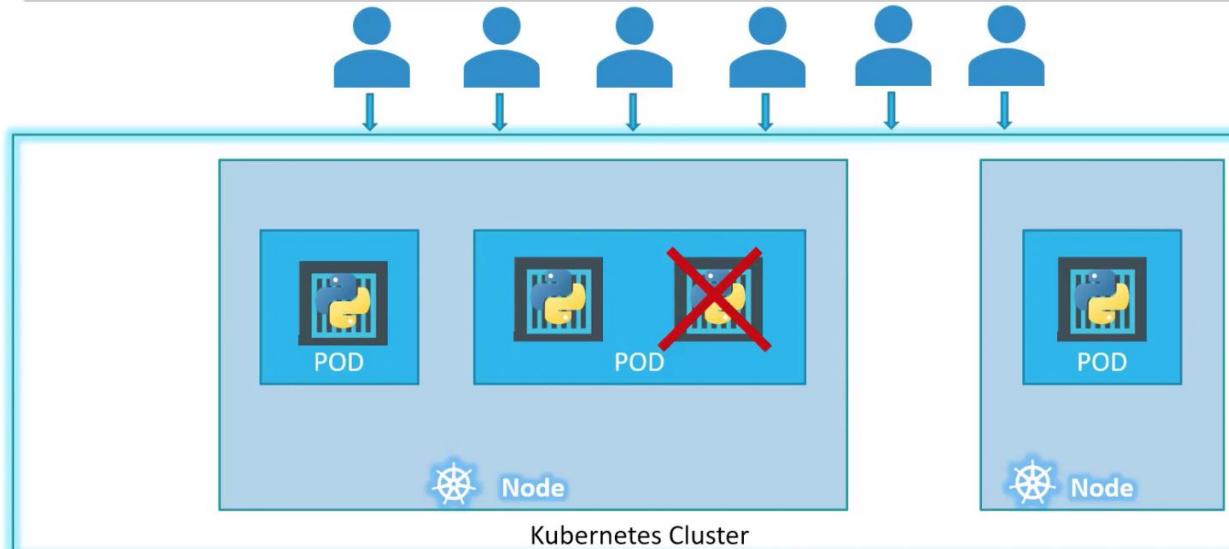
POD



Inside of a Pod, we can have multiple containers running, but this is not recommended as it is,

Usually a container runs in a single Pod, and More instances are generated as user Traffic increases

POD

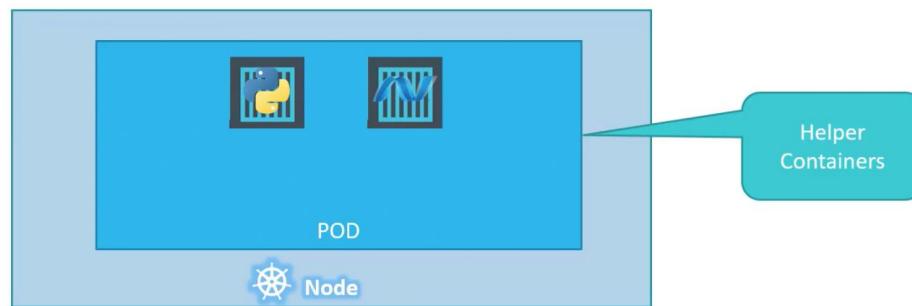


As stated Previously, we should not run two containers in a single,

But there might be cases where this is efficient, as Example is when one container, is acting as a helper (providing services to the main application).

It might be better to have both running in the same Pod

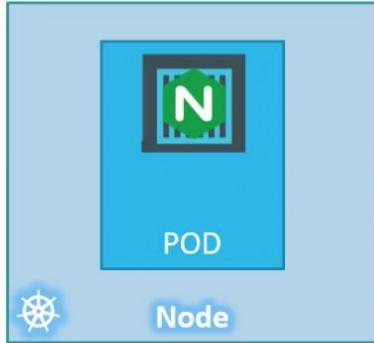
Multi-Container PODs



Simple Example running a pod in kubernetes

kubectl

```
kubectl run nginx --image nginx
```



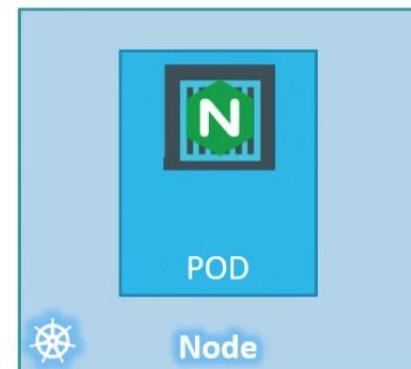
To see the list of pod in a particular namespace type

Kubectl get Pods

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	ContainerCreating	0	6s

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	34s



ReplicaSets

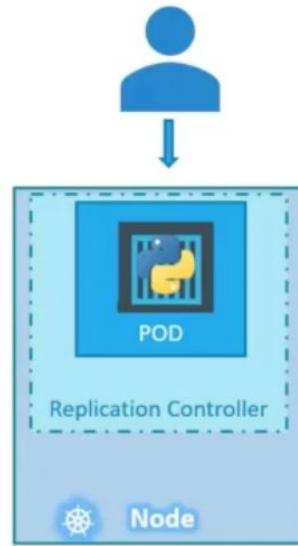
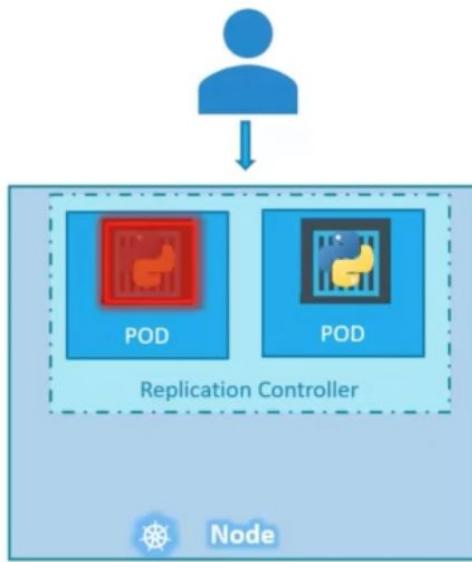
Replication Controller



Replica sets are simply copies of our pods and that keep the desired state of running Pods in the cluster to be the actual state of the pods

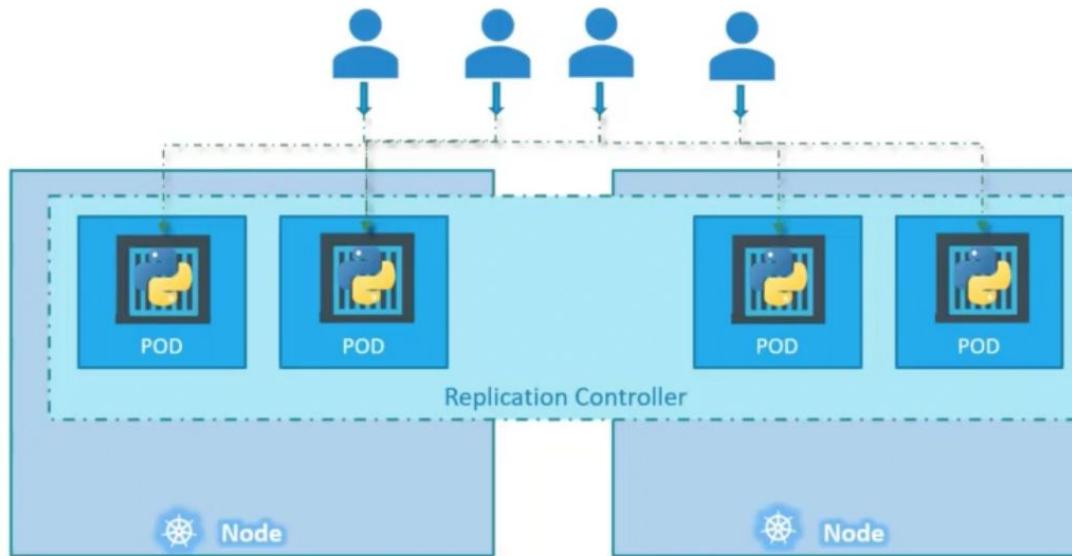
This is important because the Pods are replaced with one is down or unhealthy

High Availability

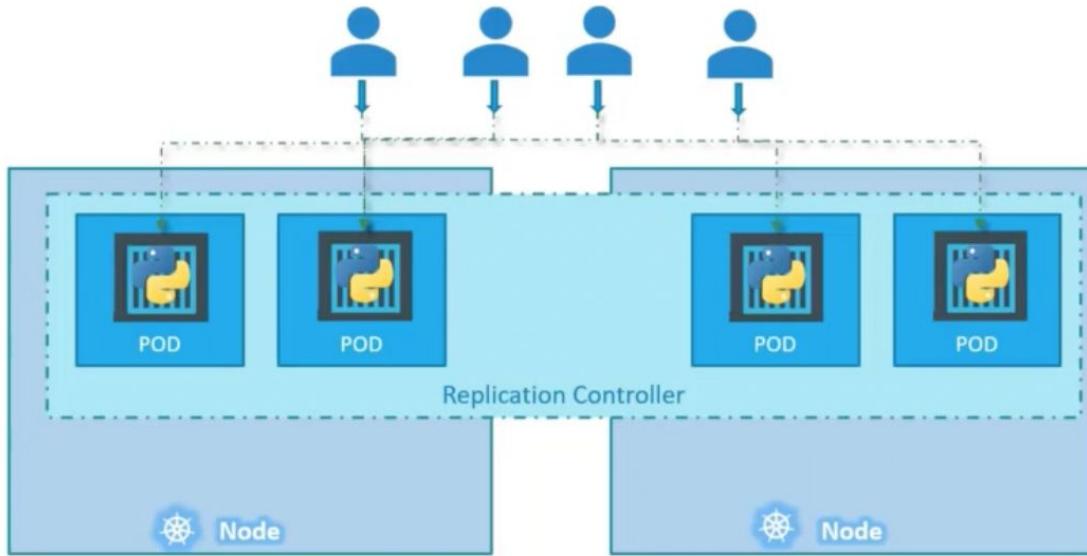


This is also Important for scaling and Load balancing when Traffic increases

Load Balancing & Scaling



Load Balancing & Scaling



Replication Controller

Replica Set

In the simplest Terms

A Replication Controller is an Older way to create Replicas

And a replica set is a newer way to do the same thing

How to create a Replication controller

We do this with a yaml, configuration file and in there we state the details of the pod to be replicated

rc-definition.yml

```
apiVersion: v1
kind: ReplicationController
metadata: Replication Controller
  name: myapp-rc
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata: POD
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
```

pod-definition.yml

```
apiVersion: v1
kind: Pod
```

Underneath the Specs for the Pod, Add the number of desired replicas, that you want

```
spec:  
  template:  
  
    metadata:  
      name: myapp-pod  
      labels:  
        app: myapp  
        type: front-end  
    spec:  
      containers:  
        - name: nginx-container  
          image: nginx  
  
    replicas: 3
```

To create a Repica set run the following command

```
> kubectl create -f rc-definition.yml  
replicationcontroller "myapp-rc" created
```

```
> kubectl get replicationcontroller
```

NAME	DESIRED	CURRENT	READY	AGE
myapp-rc	3	3	3	19s

Finally run

Kubectl get pods

To get all the replicas

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-rc-41vk9	1/1	Running	0	20s
myapp-rc-mc2mf	1/1	Running	0	20s
myapp-rc-px9pz	1/1	Running	0	20s

How to create a Replica set

We still do this with a yaml, configuration file and in there we state the details of the pod to be replicated,

But with some minor differences from a replication Controller

replicaset-definition.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
```

pod-definition.yml

```
apiVersion: v1
kind: Pod
```

Some changes you should notice would be from the

- apiVersion
 - This is set to apps/v1
- Kind
 - ReplicaSet
- Selector: this is to state the pods in this replica set

```
: selector:  
  matchLabels:  
    type: front-end
```

To create a replica set from a yam file

Run the following.

```
> kubectl create -f replicaset-definition.yml
```

```
replicaset "myapp-replicaset" created
```

```
> kubectl get replicaset
```

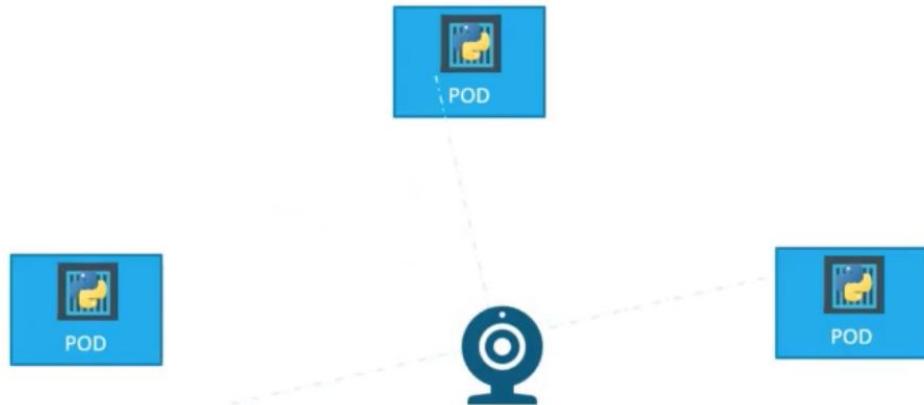
NAME	DESIRED	CURRENT	READY	AGE
myapp-replicaset	3	3	3	19s

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-replicaset-9ddl9	1/1	Running	0	45s
myapp-replicaset-9jtpx	1/1	Running	0	45s
myapp-replicaset-hq84m	1/1	Running	0	45s

Why Do we need Replica sets

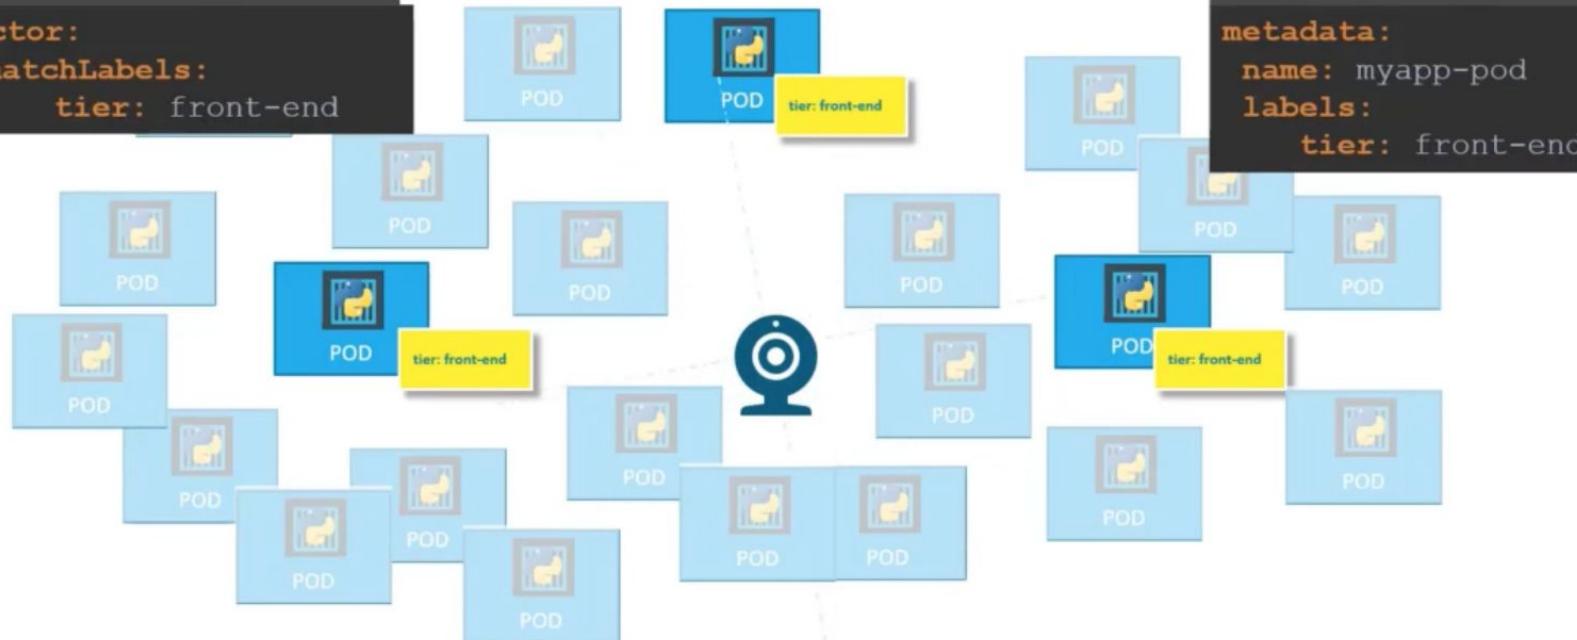
Labels and Selectors



Labels and Selectors

```
replicaset-definition.yml
```

```
selector:  
  matchLabels:  
    tier: front-end
```



```
pod-definition.yml
```

```
metadata:  
  name: myapp-pod  
  labels:  
    tier: front-end
```

Scale

```
> kubectl replace -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 replicaset myapp-replicaset
```



```
replicaset-definition.yml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 6
  selector:
    matchLabels:
      type: front-end
```

commands

```
> kubectl create -f replicaset-definition.yml
```

```
> kubectl get replicaset
```

```
> kubectl delete replicaset myapp-replicaset
```

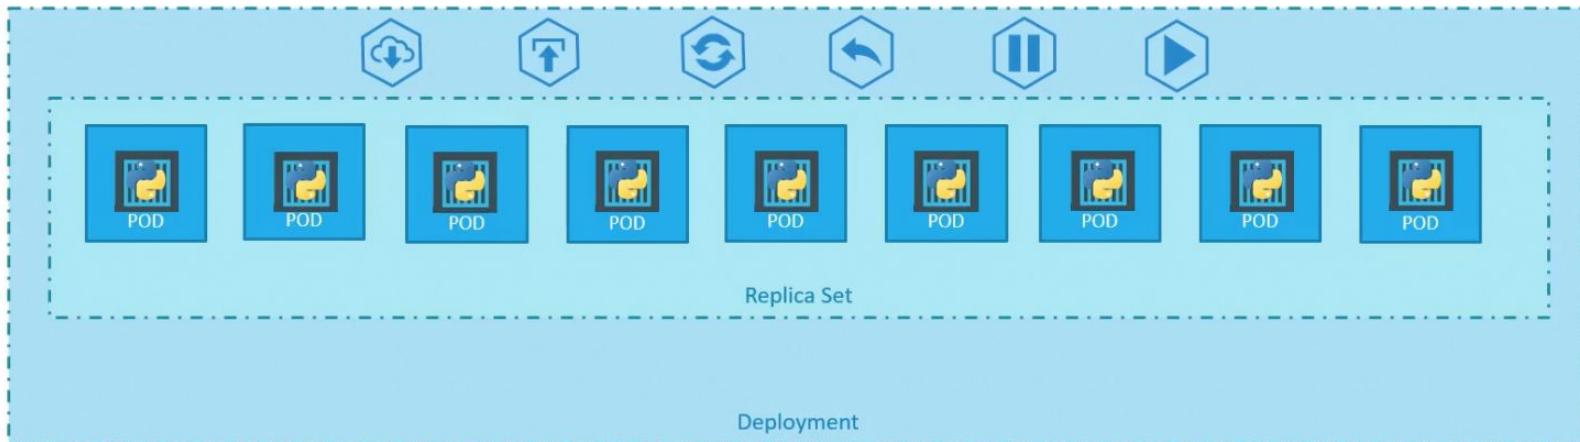
*Also deletes all underlying PODs

```
> kubectl replace -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 -f replicaset-defi
```


Deployments

Deployment



Definition

```
deployment-definition.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
      spec:
        containers:
        - name: nginx-container
          image: nginx
replicas: 3
selector:
  matchLabels:
    type: front-end
```

Pause



commands

```
> kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/myapp-deployment	3	3	3	3	9h
NAME	DESIRED	CURRENT	READY	AGE	
rs/myapp-deployment-6795844b58	3	3	3	9h	
NAME	READY	STATUS	RESTARTS	AGE	
po/myapp-deployment-6795844b58-5rbj1	1/1	Running	0	9h	
po/myapp-deployment-6795844b58-h4w55	1/1	Running	0	9h	
po/myapp-deployment-6795844b58-1fjhv	1/1	Running	0	9h	

Certification Tip!

Here's a tip!

As you might have seen already, it is a bit difficult to create and edit YAML files. Especially in the CLI. During the exam, you might find it difficult to copy and paste YAML files from browser to terminal. Using the `kubectl run` command can help in generating a YAML template. And sometimes, you can even get away with just the `kubectl run` command without having to create a YAML file at all. For example, if you were asked to create a pod or deployment with specific name and image you can simply run the `kubectl run` command.

Use the below set of commands and try the previous practice tests again, but this time try to use the below commands instead of YAML files. Try to use these as much as you can going forward in all exercises

Reference (Bookmark this page for exam. It will be very handy):

<https://kubernetes.io/docs/reference/kubectl/conventions/>

Create an NGINX Pod

```
kubectl run nginx --image=nginx
```

Generate POD Manifest YAML file (-o yaml). Don't create it(-dry-run)

```
kubectl run nginx --image=nginx --dry-run=client -o yaml
```

Create a deployment

```
kubectl create deployment --image=nginx nginx
```

Generate Deployment YAML file (-o yaml). Don't create it(-dry-run)

```
kubectl create deployment --image=nginx nginx --dry-run=client -o yaml
```

Generate Deployment YAML file (-o yaml). Don't create it(-dry-run) with 4 Replicas (-replicas=4)

```
kubectl create deployment --image=nginx nginx --dry-run=client -o yaml >  
nginx-deployment.yaml
```

Save it to a file, make necessary changes to the file (for example, adding more replicas) and then create the deployment.

```
kubectl create -f nginx-deployment.yaml
```

OR

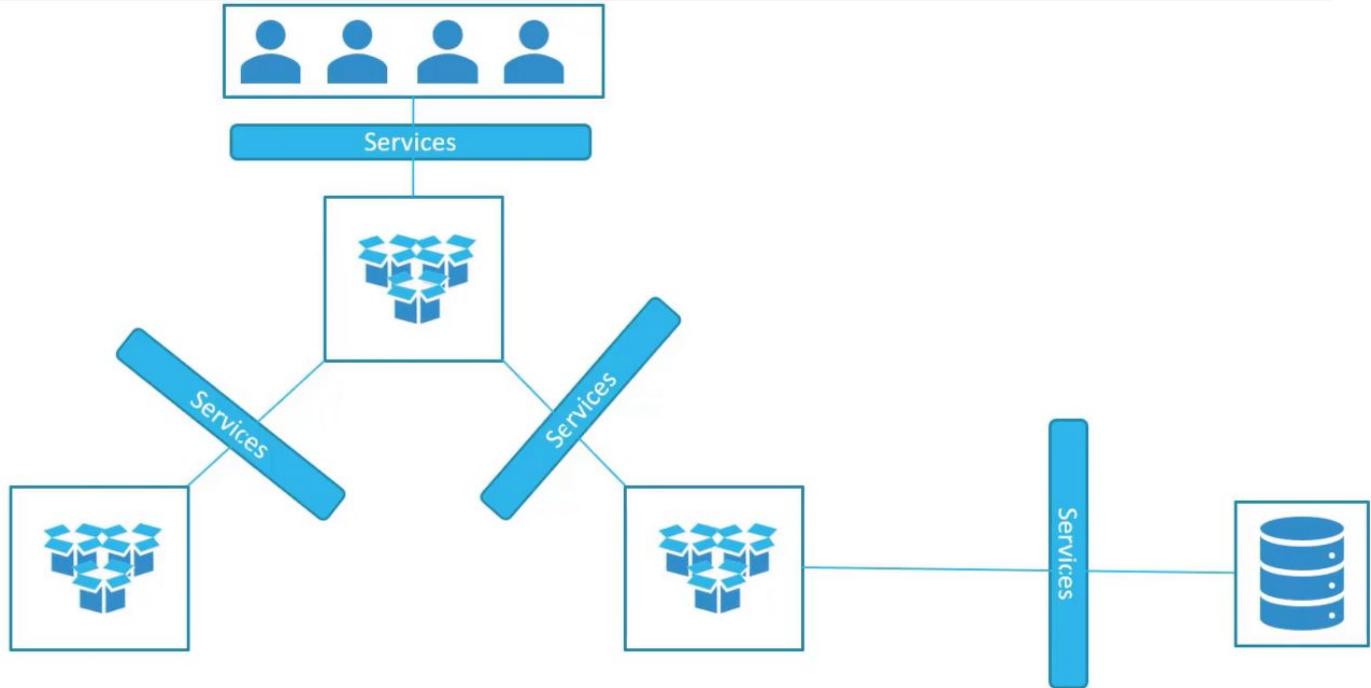
In k8s version 1.19+, we can specify the --replicas option to create a deployment with 4 replicas.

```
kubectl create deployment --image=nginx nginx --replicas=4 --dry-run=client -o yaml > nginx-deployment.yaml
```

SERVICES

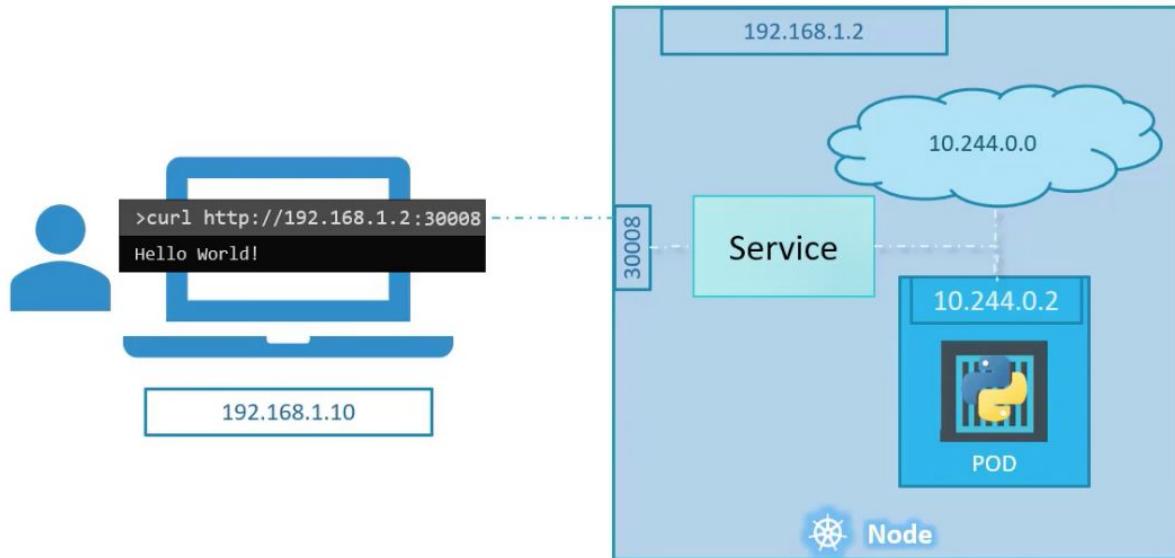


Services



Services Makes available an Internal network that, users can interact with the Pods and a Pod (Application) can connect To another Pod (Application).

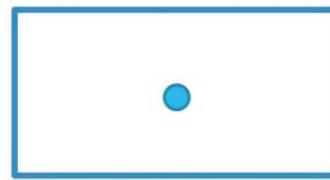
Service



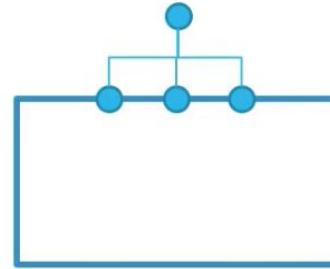
Services Types



NodePort



ClusterIP



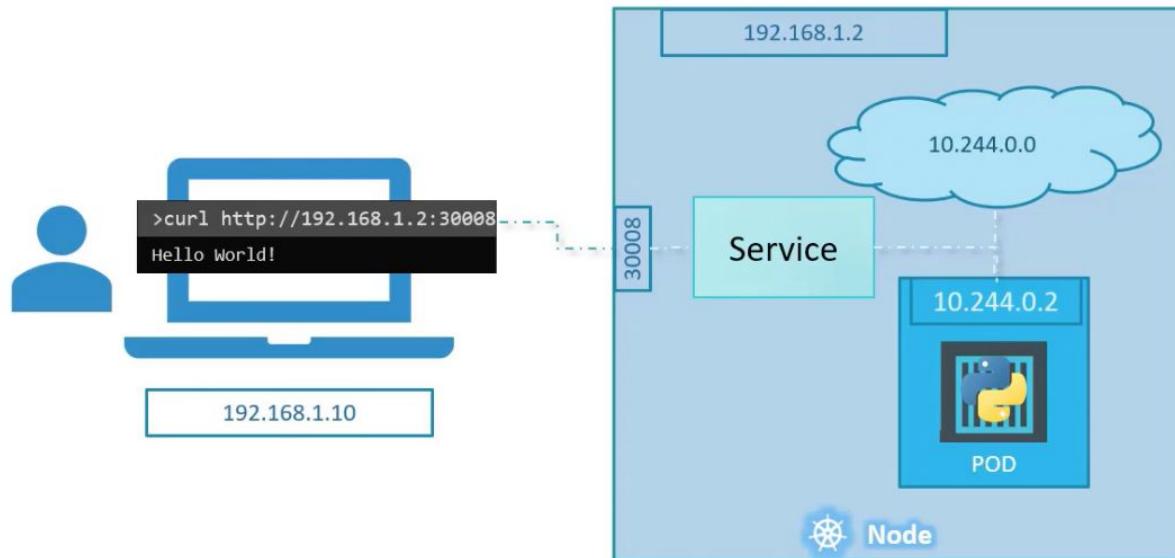
LoadBalancer

NodePort: This Maps a Port on the Pod to a Port on the containers, so external users can access the Application via the node port IP

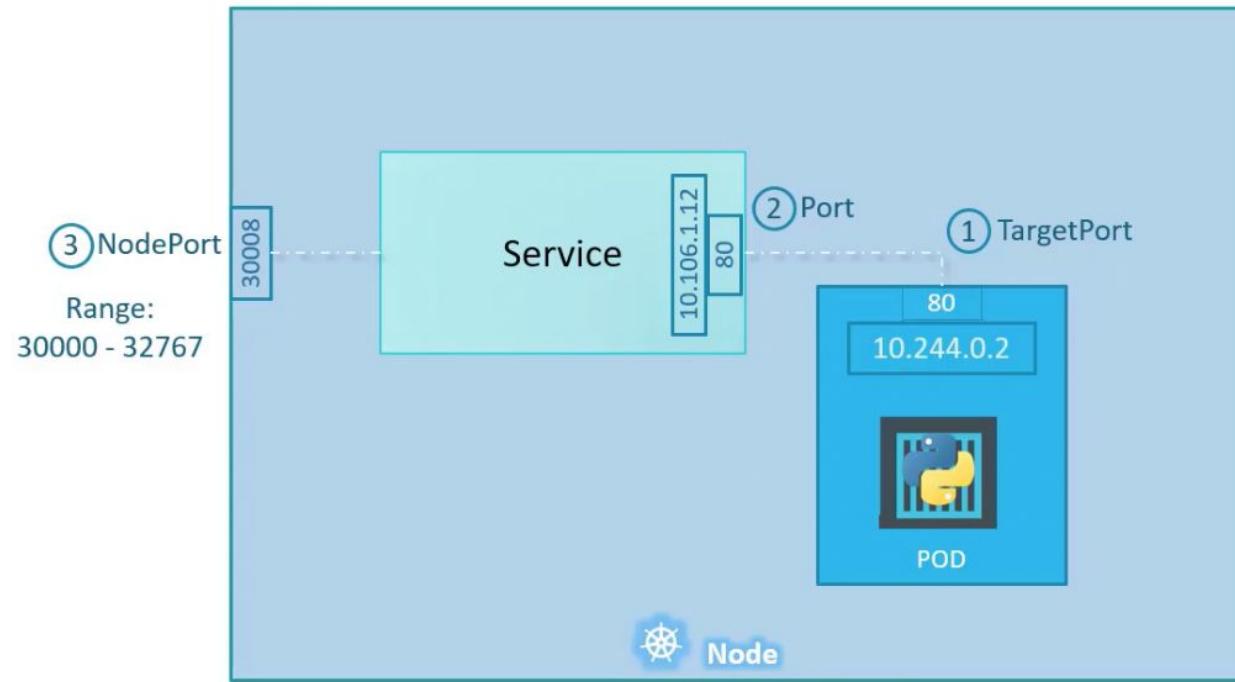
NodePort



Service - NodePort



Service - NodePort



service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service

spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

Creating a Service

```
Kubectl create -f service-defination.yaml
```

Get the list of services

```
Kubectl get services
```

```
Kubectl get svc
```

Adding Selectors And Labels

Service - NodePort

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service

spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
```

```
pod-definition.yml
```

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

Add labels from Pod

Service - NodePort

`service-definition.yml`

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: myapp
    type: front-end
```

`pod-definition.yml`

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
spec:
  containers:
    - name: nginx-container
      image: nginx
```

```
> kubectl create -f service-definition.yml
```

```
service "myapp-service" created
```

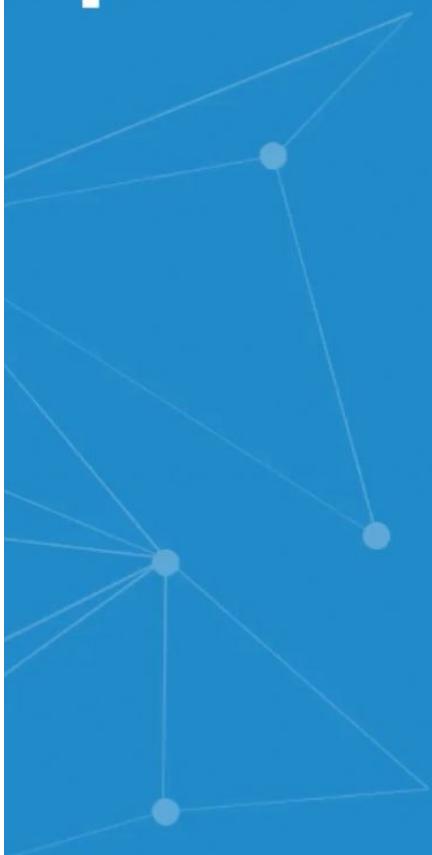
```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
myapp-service	NodePort	10.106.127.123	<none>	80:30008/TCP	5m

```
> curl http://192.168.1.2:30008
```

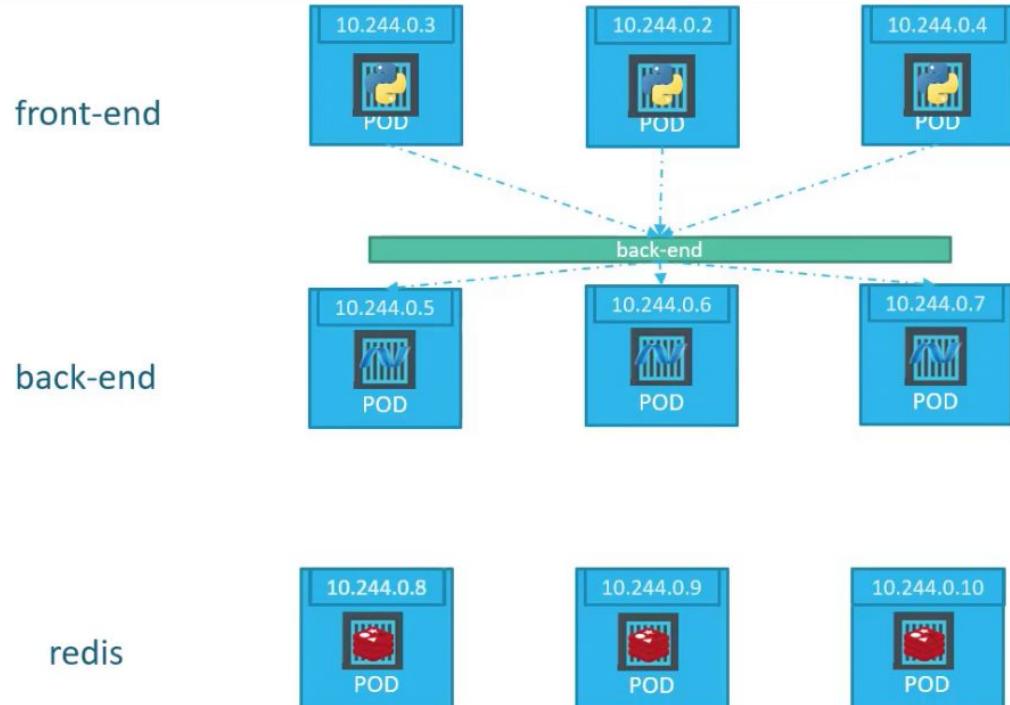
```
<html>
<head>
<title>Welcome to nginx!</title>
<style>
    body {
        width: 35em;
        margin: 0 auto;
        font-family: Tahoma, Verdana, Arial, sans-serif;
    }
</style>
</head>
<body>
```

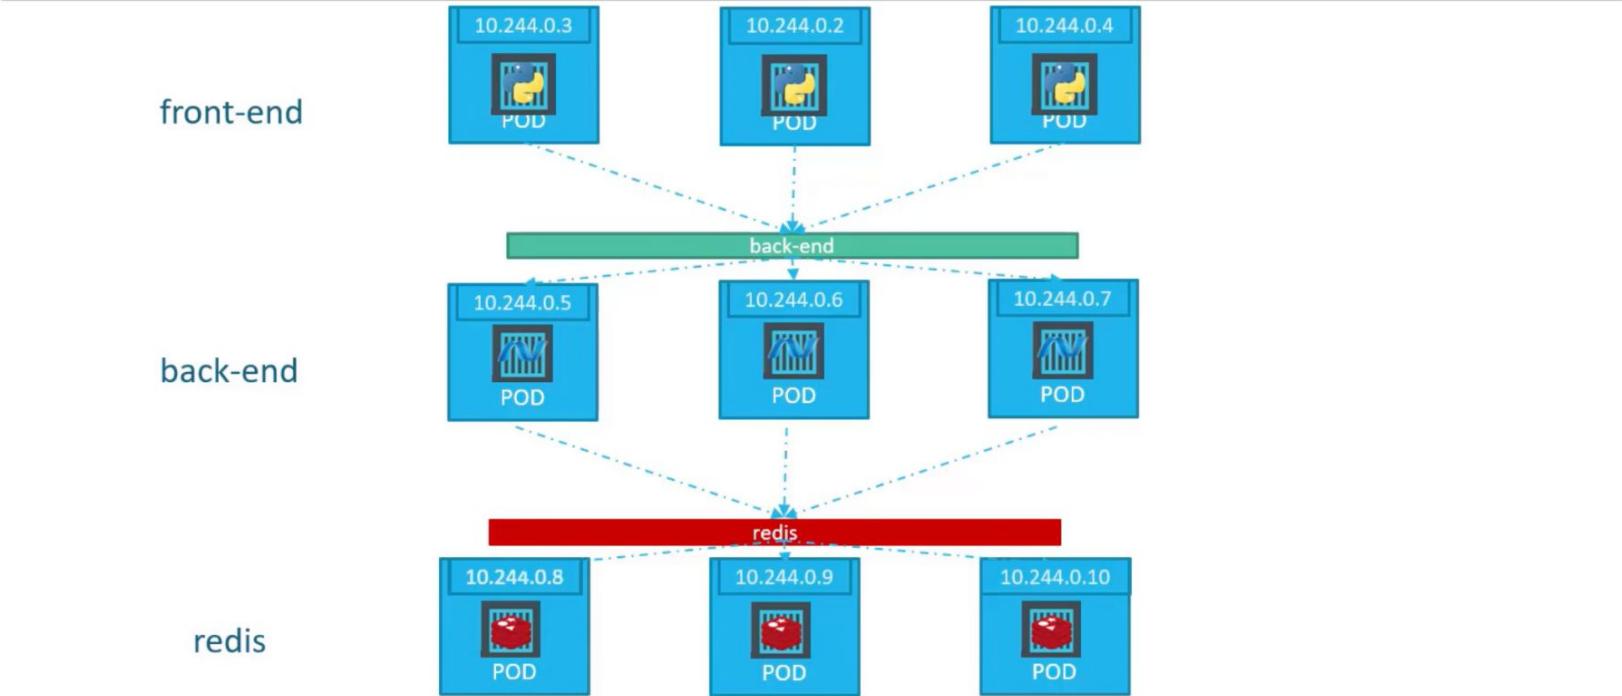
CLUSTER IP

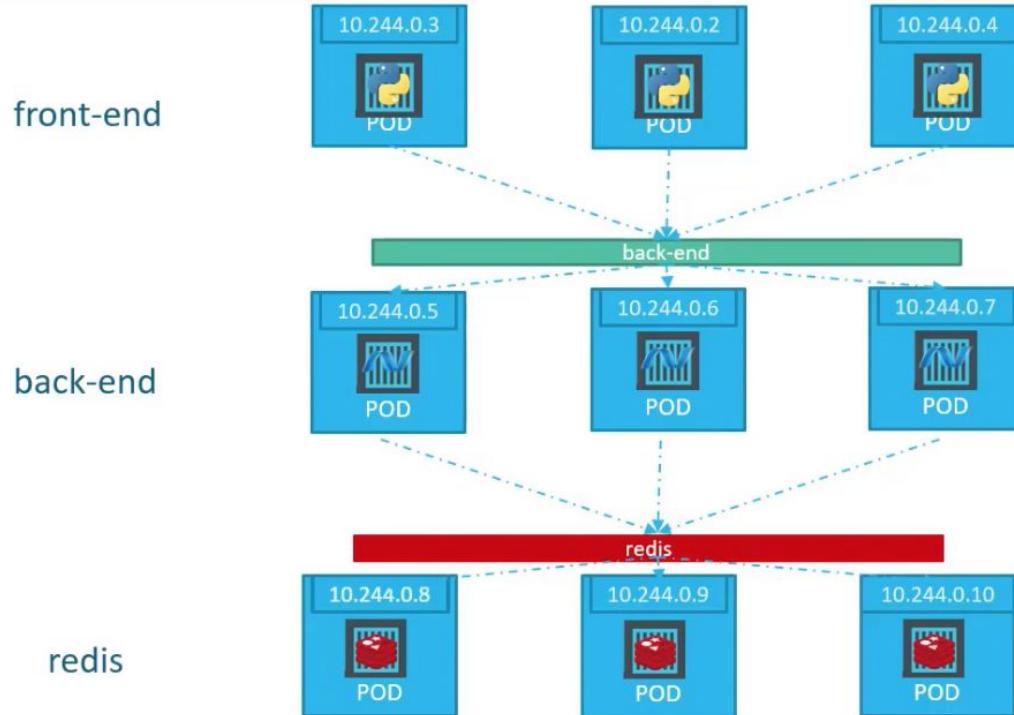


In a multi-tier architecture where we have multiple applications on Multiple Pods, trying to connect to each other over the network, a service Cluster IP is an ideal approach to handle networking.

A kubernetes Cluster IP Service, can help us to group the
Pods together, so we can access them in a Group







Creating A Cluster IP Service

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: back-end

spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80

  selector:
```

```
pod-definition.yml
```

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: back-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: back-end

spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80

  selector:
    app: myapp
    type: back-end
```

pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:

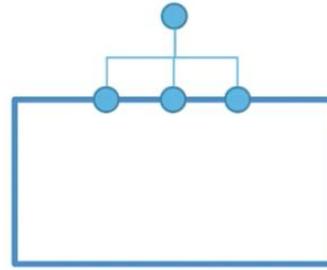
spec:
  containers:
    - name: nginx-container
      image: nginx
```

```
> kubectl create -f service-definition.yml  
service "back-end" created
```

```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
back-end	ClusterIP	10.106.127.123	<none>	80/TCP	2m

Service- Load Balancer



Services – Loadbalancer

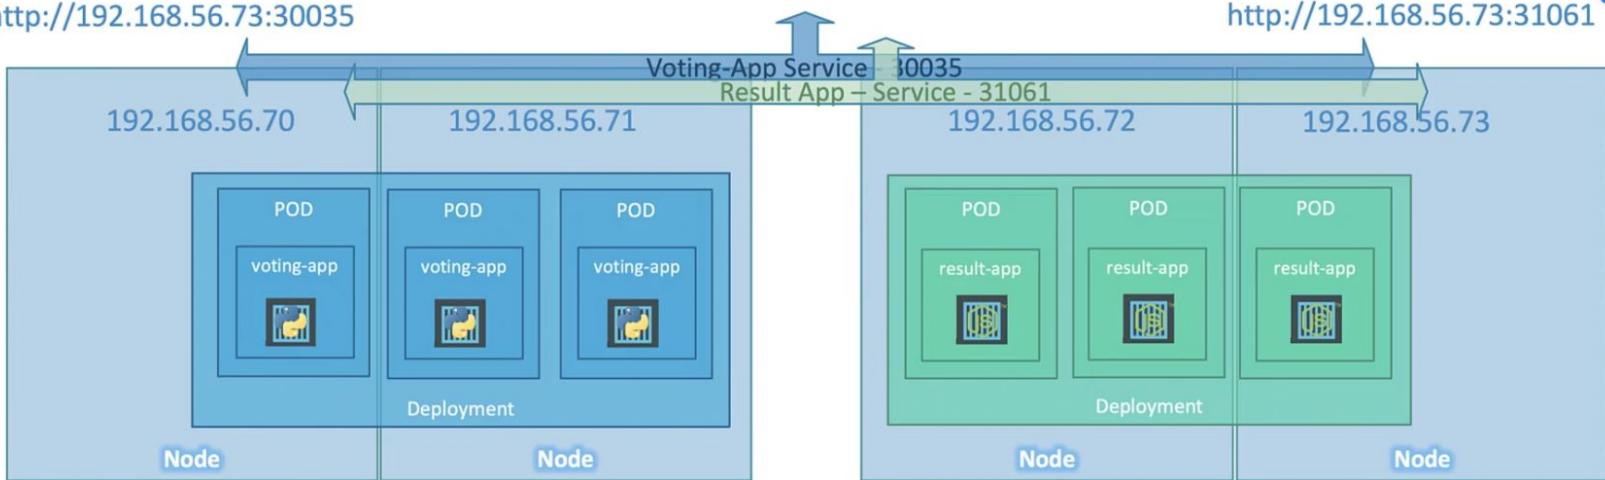
Example voting app



<http://192.168.56.70:30035>
<http://192.168.56.71:30035>
<http://192.168.56.72:30035>
<http://192.168.56.73:30035>

<http://example-vote.com>
<http://example-result.com>

<http://192.168.56.70:31061>
<http://192.168.56.71:31061>
<http://192.168.56.72:31061>
<http://192.168.56.73:31061>

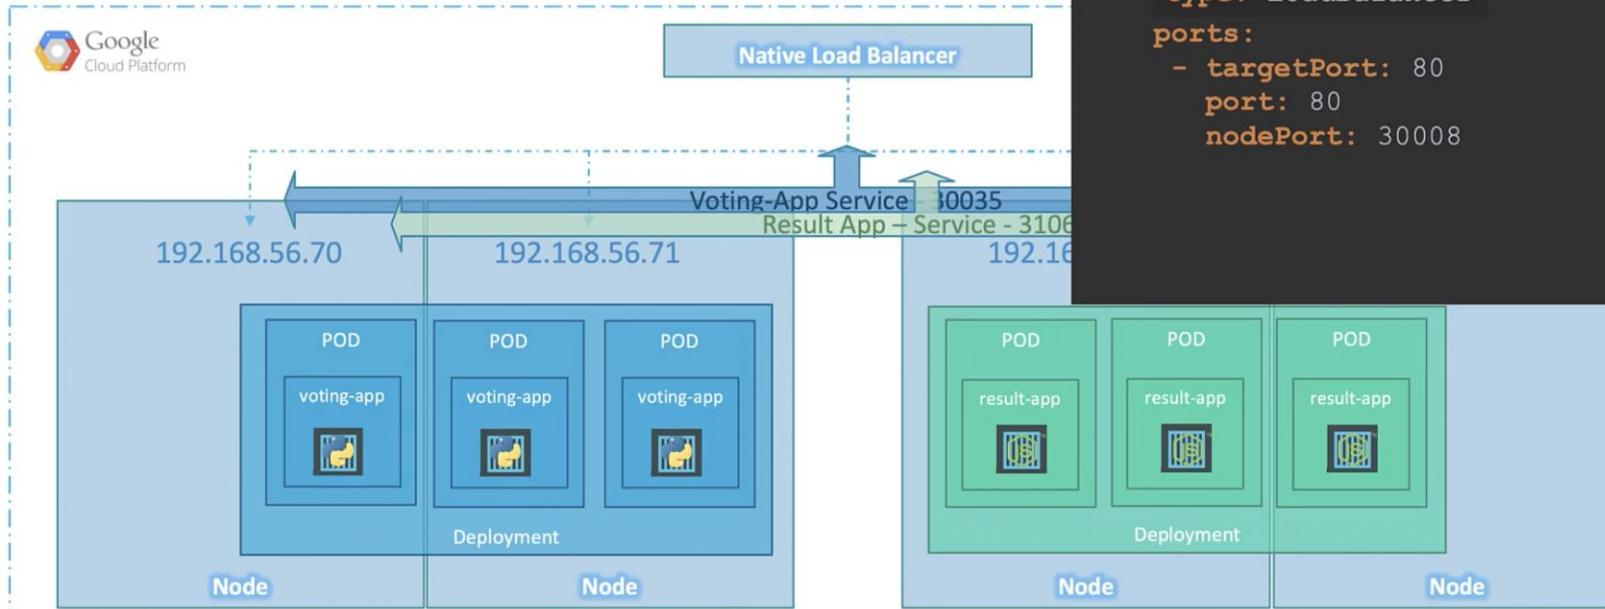


Services – Loadbalancer

Example voting app



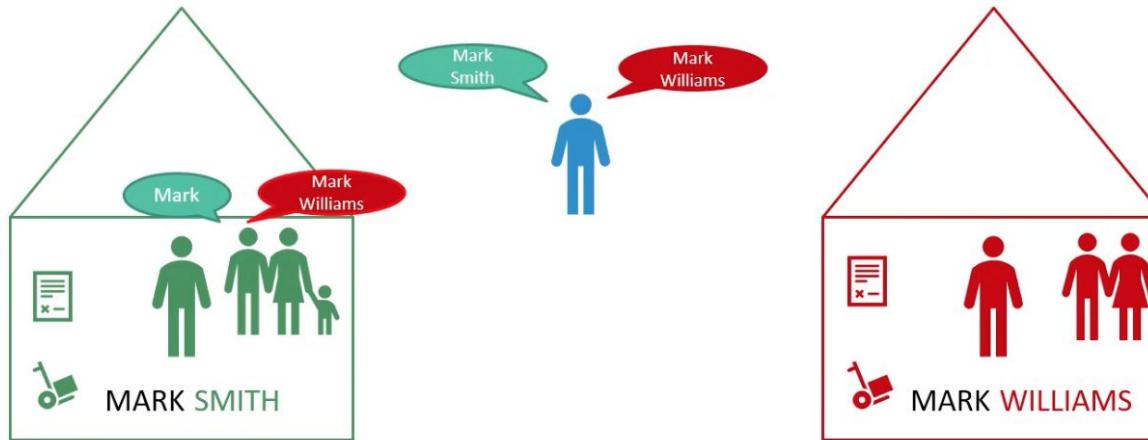
<http://example-vote.com>
<http://example-result.com>



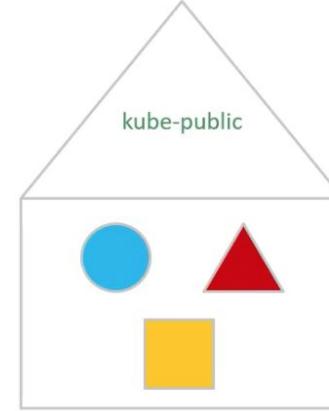
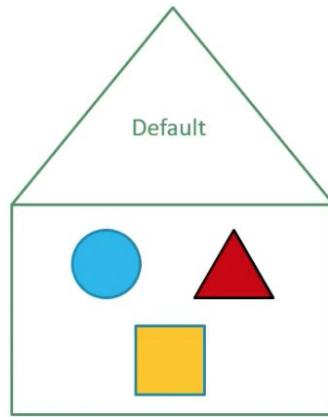
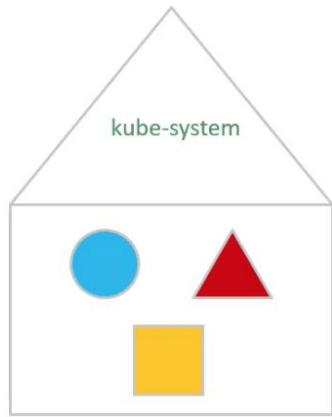
Namespaces



Namespaces

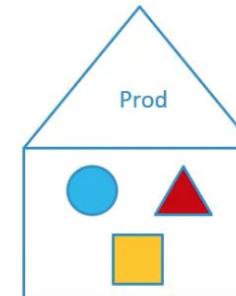
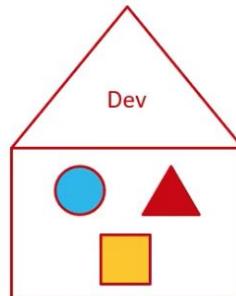
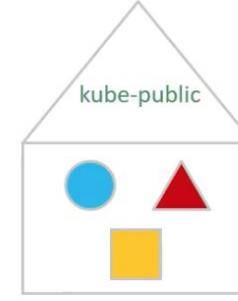
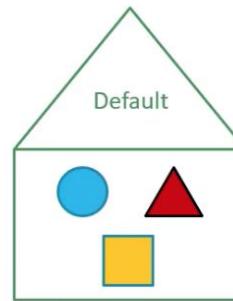
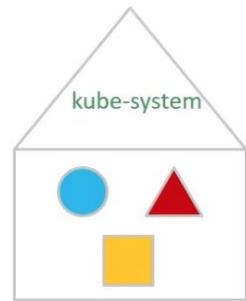


Namespaces



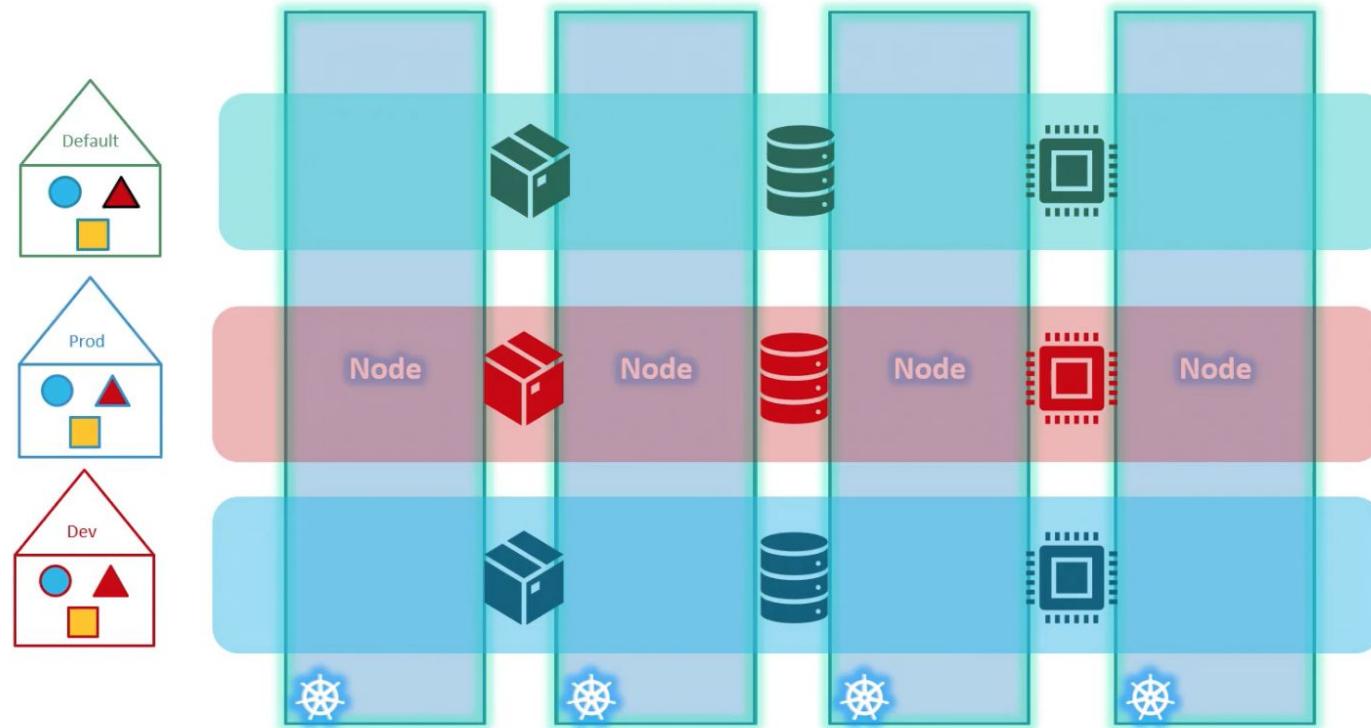
Namespaces

Namespace - Isolation



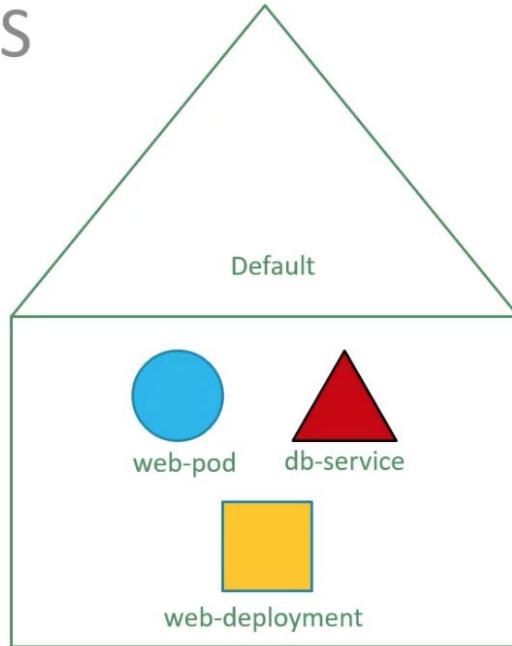
Namespaces

Namespace – Resource Limits

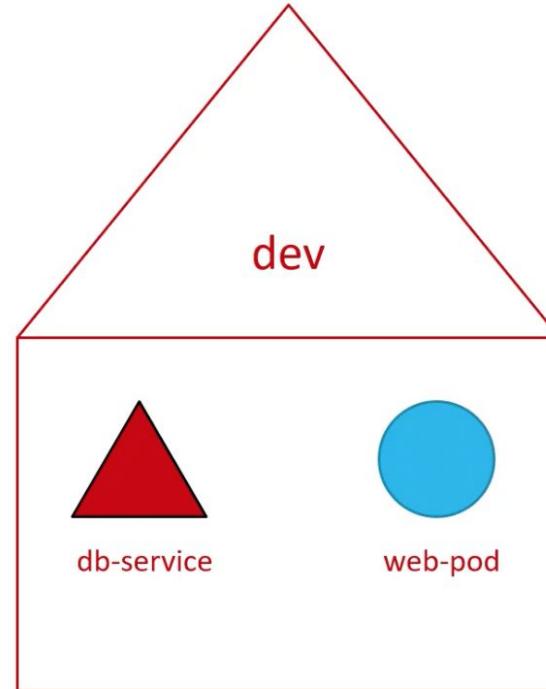


Namespaces

DNS



```
mysql.connect("db-service")
```



```
mysql.connect("db-service.dev.svc.cluster.local")
```

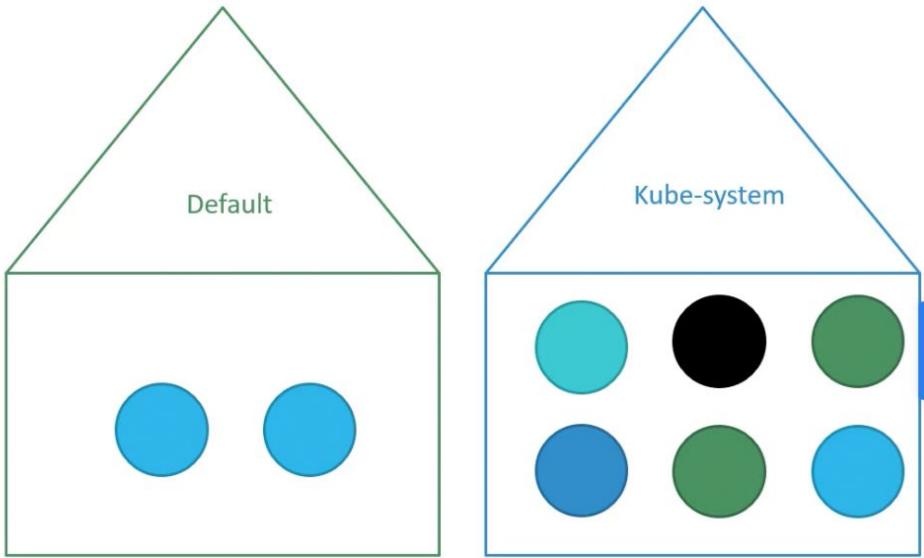
Namespaces

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
Pod-1	1/1	Running	0	3d
Pod-2	1/1	Running	0	3d

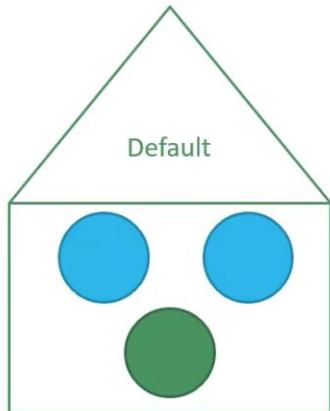
```
> kubectl get pods --namespace=kube-system
```

NAME	READY	STATUS	RESTAR
coredns-78fcdf6894-92d52	1/1	Running	7
coredns-78fcdf6894-jx25g	1/1	Running	7
etcd-master	1/1	Running	7
kube-apiserver-master	1/1	Running	7
kube-controller-manager-master	1/1	Running	7
kube-flannel-ds-amd64-hz4cf	1/1	Running	14
kube-proxy-4b8tn	1/1	Running	7
kube-proxy-98db4	1/1	Running	7
kube-proxy-jjrbs	1/1	Running	7
kube-scheduler-master	1/1	Running	7



Namespaces

```
> kubectl create -f pod-definition.yml  
pod/myapp-pod created
```



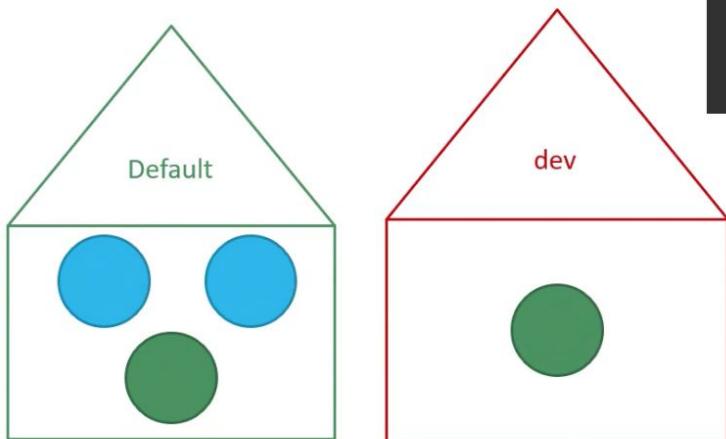
```
pod-definition.yml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: myapp-pod  
  labels:  
    app: myapp  
    type: front-end  
spec:  
  containers:  
  - name: nginx-container  
    image: nginx
```

Namespaces

```
> kubectl create -f pod-definition.yml  
pod/myapp-pod created
```

```
> kubectl create -f pod-definition.yml --namespace=dev  
pod/myapp-pod created
```

```
pod-definition.yml  
  
apiVersion: v1  
kind: Pod  
  
metadata:  
  name: myapp-pod  
  labels:  
    app: myapp  
    type: front-end  
  
spec:  
  containers:  
  - name: nginx-container  
    image: nginx
```

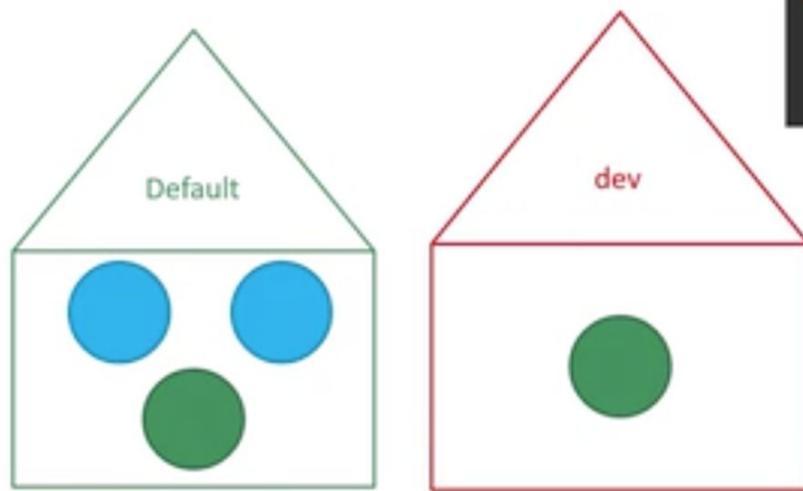


```
> kubectl create -f pod-definition.yml  
pod/myapp-pod created
```

```
> kubectl create -f pod-definition.yml  
pod/myapp-pod created
```

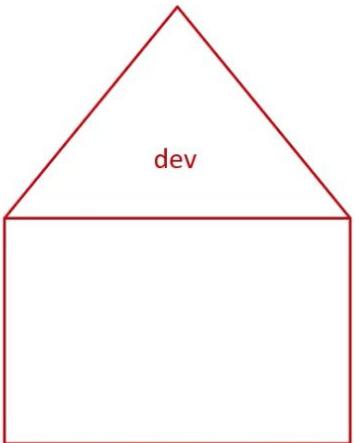
```
pod-definition.yml
```

```
apiVersion: v1  
kind: Pod  
  
metadata:  
  name: myapp-pod  
  namespace: dev  
  
labels:  
  app: myapp  
  type: front-end  
  
spec:  
  containers:  
    - name: nginx-container  
      image: nginx
```



Namespaces

Create Namespace



```
namespace-dev.yml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: dev
```

```
> kubectl create -f namespace-dev.yml
```

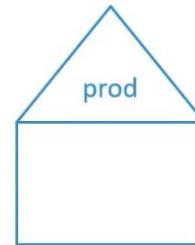
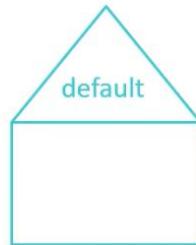
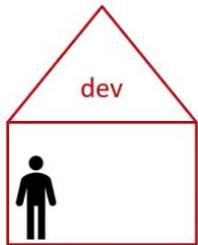
```
namespace/dev created
```

```
> kubectl create namespace dev
```

```
namespace/dev created
```

Namespaces

Switch



```
> kubectl get pods --namespace=dev
```

```
> kubectl get pods
```

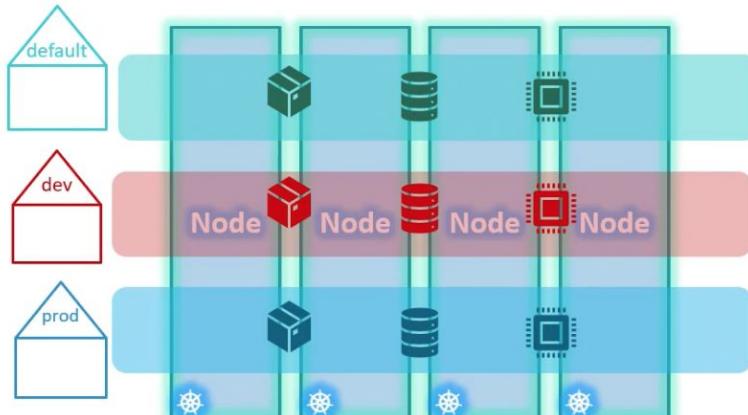
```
> kubectl get pods --namespace=prod
```

```
> kubectl config set-context $(kubectl config current-context) --namespace=dev
```

```
> kubectl get pod
```

Namespaces

Resource Quota



```
Compute-quota.yaml
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-quota
  namespace: dev
spec:
  hard:
    pods: "10"
    requests.cpu: "4"
    requests.memory: 5Gi
    limits.cpu: "10"
    limits.memory: 10Gi
```

```
> kubectl create -f compute-quota.yaml
```

Imperative and Declarative

Imperative
vs
Declarative

Imperative vs Declarative

Imperative



Street A

1. Right

Street B

2. Left

4. Right

Street D

3. Left



Declarative



Street A

Street B

Street C

Street D

1. Go to Tom's house



Imperative vs Declarative

Infrastructure as Code

Imperative

1. Provision a VM by the name ‘web-server’
2. Install NGINX Software on it
3. Edit configuration file to use port ‘8080’
4. Edit configuration file to web path ‘/var/www/nginx’
5. Load web pages to ‘/var/www/nginx’ from GIT Repo - X
6. Start NGINX server

Declarative

VM Name: web-server
Package: nginx
Port: 8080
Path: /var/www/nginx
Code: GIT Repo - X

Imperative vs Declarative

Kubernetes

Imperative

```
> kubectl run --image=nginx nginx
```

```
> kubectl create deployment --image=nginx nginx
```

```
> kubectl expose deployment nginx --port 80
```

```
> kubectl edit deployment nginx
```

```
> kubectl scale deployment nginx --replicas=5
```

```
> kubectl set image deployment nginx nginx=nginx:1.18
```

```
> kubectl create -f nginx.yaml
```

```
> kubectl replace -f nginx.yaml
```

```
> kubectl delete -f nginx.yaml
```

Declarative

Declarative

```
> kubectl apply -f nginx.yaml
```

Imperative vs Declarative

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

```
> kubectl replace -f nginx.yaml
```

```
> kubectl replace --force -f nginx.yaml
```

```
> kubectl create -f nginx.yaml
```

```
Error from server (AlreadyExists): error when creating "nginx.yaml": pods "myapp-pod" already exists
```

```
> kubectl replace -f nginx.yaml
```

```
Error from server (Conflict): error when replacing "nginx.yaml": Operation cannot be fulfilled on pods "myapp-pod"
```

```
nginx.yaml
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx:1.18
```

Imperative vs Declarative

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

nginx.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
  - name: nginx-container
    image: nginx
```

pod-definition

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
  - name: nginx-container
    image: nginx

status:
  conditions:
  - lastProbeTime: null
    status: "True"
    type: Initialized
```

Imperative vs Declarative

Imperative Object Configuration Files

Create Objects

```
> kubectl create -f nginx.yaml
```

Update Objects

```
> kubectl edit deployment nginx
```

```
> kubectl replace -f nginx.yaml
```

```
> kubectl replace --force -f nginx.yaml
```

```
> kubectl create -f nginx.yaml
```

```
Error from server (AlreadyExists): error when creating "nginx.yaml": pods "myapp-pod" already exists
```

```
> kubectl replace -f nginx.yaml
```

```
Error from server (Conflict): error when replacing "nginx.yaml": Operation cannot be fulfilled on pods  
"myapp-pod"
```

```
nginx.yaml
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
    - name: nginx-container
      image: nginx:1.18
```

Imperative vs Declarative

Declarative

Create Objects

```
> kubectl apply -f nginx.yaml
```

```
> kubectl apply -f /path/to/config-files
```

Update Objects

```
> kubectl apply -f nginx.yaml
```

```
nginx.yaml
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
    - name: nginx-container
      image: nginx:1.18
```

Certification Tips – Imperative Commands with Kubectl

While you would be working mostly the declarative way – using definition files, imperative commands can help in getting one time tasks done quickly, as well as generate a definition template easily. This would help save considerable amount of time during your exams.

Before we begin, familiarize with the two options that can come in handy while working with the below commands:

- `--dry-run`: By default as soon as the command is run, the resource will be created. If you simply want to test your command, use the `--dry-run=client` option. This will not create the resource, instead, tell you whether the resource can be created and if your command is right.
- o `yaml`: This will output the resource definition in YAML format on screen.

Use the above two in combination to generate a resource definition file quickly, that you can then modify and create resources as required, instead of creating the files from scratch.

POD

Create an NGINX Pod

```
kubectl run nginx --image=nginx
```

Generate POD Manifest YAML file (-o yaml). Don't create it(--dry-run)

```
kubectl run nginx --image=nginx --dry-run=client -o yaml
```

Deployment

Create a deployment

```
kubectl create deployment --image=nginx nginx
```

Generate Deployment YAML file (-o yaml). Don't create it(-dry-run)

```
kubectl create deployment --image=nginx nginx --dry-run=client -o yaml
```

Generate Deployment with 4 Replicas

```
kubectl create deployment nginx --image=nginx --replicas=4
```

You can also scale a deployment using the `kubectl scale` command.

```
kubectl scale deployment nginx --replicas=4
```

Another way to do this is to save the YAML definition to a file and modify

```
kubectl create deployment nginx --image=nginx --dry-run=client -o yaml > nginx-deployment.yaml
```

You can then update the YAML file with the replicas or any other field before creating the deployment.

Service

Create a Service named redis-service of type ClusterIP to expose pod redis on port 6379

```
kubectl expose pod redis --port=6379 --name redis-service --dry-run=client -o yaml
```

(This will automatically use the pod's labels as selectors)

Or

```
kubectl create service clusterip redis --tcp=6379:6379 --dry-run=client -o yaml
```

(This will not use the pods labels as selectors, instead it will assume selectors as app=redis. **You cannot pass in selectors as an option.** So it does not work very well if your pod has a different label set. So generate the file and modify the selectors before creating the service)

Create a Service named nginx of type NodePort to expose pod nginx's port 80 on port 30080 on the nodes:

```
kubectl expose pod nginx --type=NodePort --port=80 --name=nginx-service --dry-run=client -o yaml
```

(This will automatically use the pod's labels as selectors, **but you cannot specify the node port**. You have to generate a definition file and then add the node port in manually before creating the service with the pod.)

Or

```
kubectl create service nodeport nginx --tcp=80:80 --node-port=30080 --dry-run=client -o yaml
```

(This will not use the pods labels as selectors)

Both the above commands have their own challenges. While one of it cannot accept a selector the other cannot accept a node port. I would recommend going with the `kubectl expose` command. If you need to specify a node port, generate a definition file using the same command and manually input the nodeport before creating the service.

Reference:

<https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands>

<https://kubernetes.io/docs/reference/kubectl/conventions/>



Kubectl
Apply

Kubectl Apply

Local file

```
nginx.yaml
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx:1.18
```

```
> kubectl apply -f nginx.yaml
```

Last applied Configuration

 Kubernetes

```
Live object configuration
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
  - name: nginx-container
    image: nginx:1.18

status:
  conditions:
  - lastProbeTime: null
    status: "True"
    type: Initialized
```

Kubectl Apply

Local file

nginx.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
    - name: nginx-container
      image: nginx:1.18
```

```
> kubectl apply -f nginx.yaml
```

Last applied Configuration

```
{
  "apiVersion": "v1",
  "kind": "Pod",
  "metadata": {
    "annotations": {},
    "labels": {
      "run": "myapp-pod",
      "type": "front-end-service"
    },
    "name": "myapp-pod",
  },
  "spec": {
    "containers": [
      {
        "image": "nginx:1.18",
        "name": "nginx-container"
      }
    ]
  }
}
```

 Kubernetes

Live object configuration

```
apiVersion: v1
kind: Pod

metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end-service
spec:
  containers:
    - name: nginx-container
      image: nginx:1.18

status:
  conditions:
    - lastProbeTime: null
      status: "True"
      type: Initialized
```